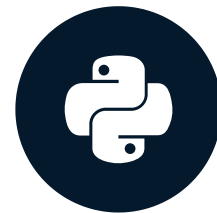


Introduction to RLHF


REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)



Mina Parham
AI Engineer


Welcome to the course!

- Instructor: Mina Parham
- Topic: Reinforcement Learning from Human Feedback (RLHF)
- AI Engineer
- Large Language Models (LLMs)
- Reinforcement Learning from Human Feedback (RLHF)




Welcome to the course!


- Instructor: Mina Parham
- Topic: Reinforcement Learning from Human Feedback (RLHF)
- AI Engineer
- Large Language Models (LLMs)
- Reinforcement Learning from Human Feedback (RLHF)




Reinforcement learning review




Reinforcement learning review




Reinforcement learning review




Reinforcement learning review



From RL to RLHF

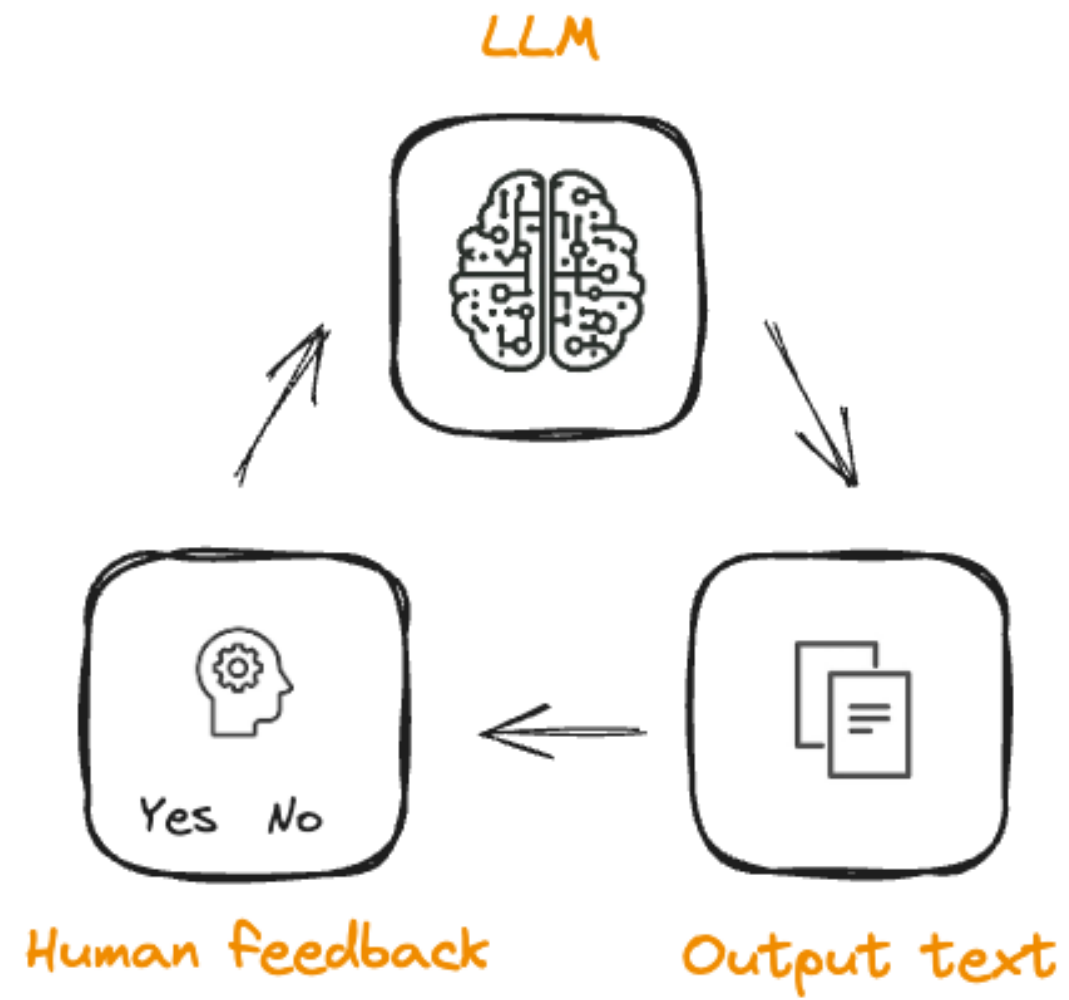


From RL to RLHF




From RL to RLHF

- Training the reward model
- Alignment with human preferences




LLM fine-tuning in RLHF

LLM



LLM fine-tuning in RLHF

- Training the initial LLM




The full RLHF process

Who wrote "Romeo and Juliet"?



Initial LLM




The full RLHF process

Who wrote "Romeo and Juliet"?




Initial LLM



A 16th Century author

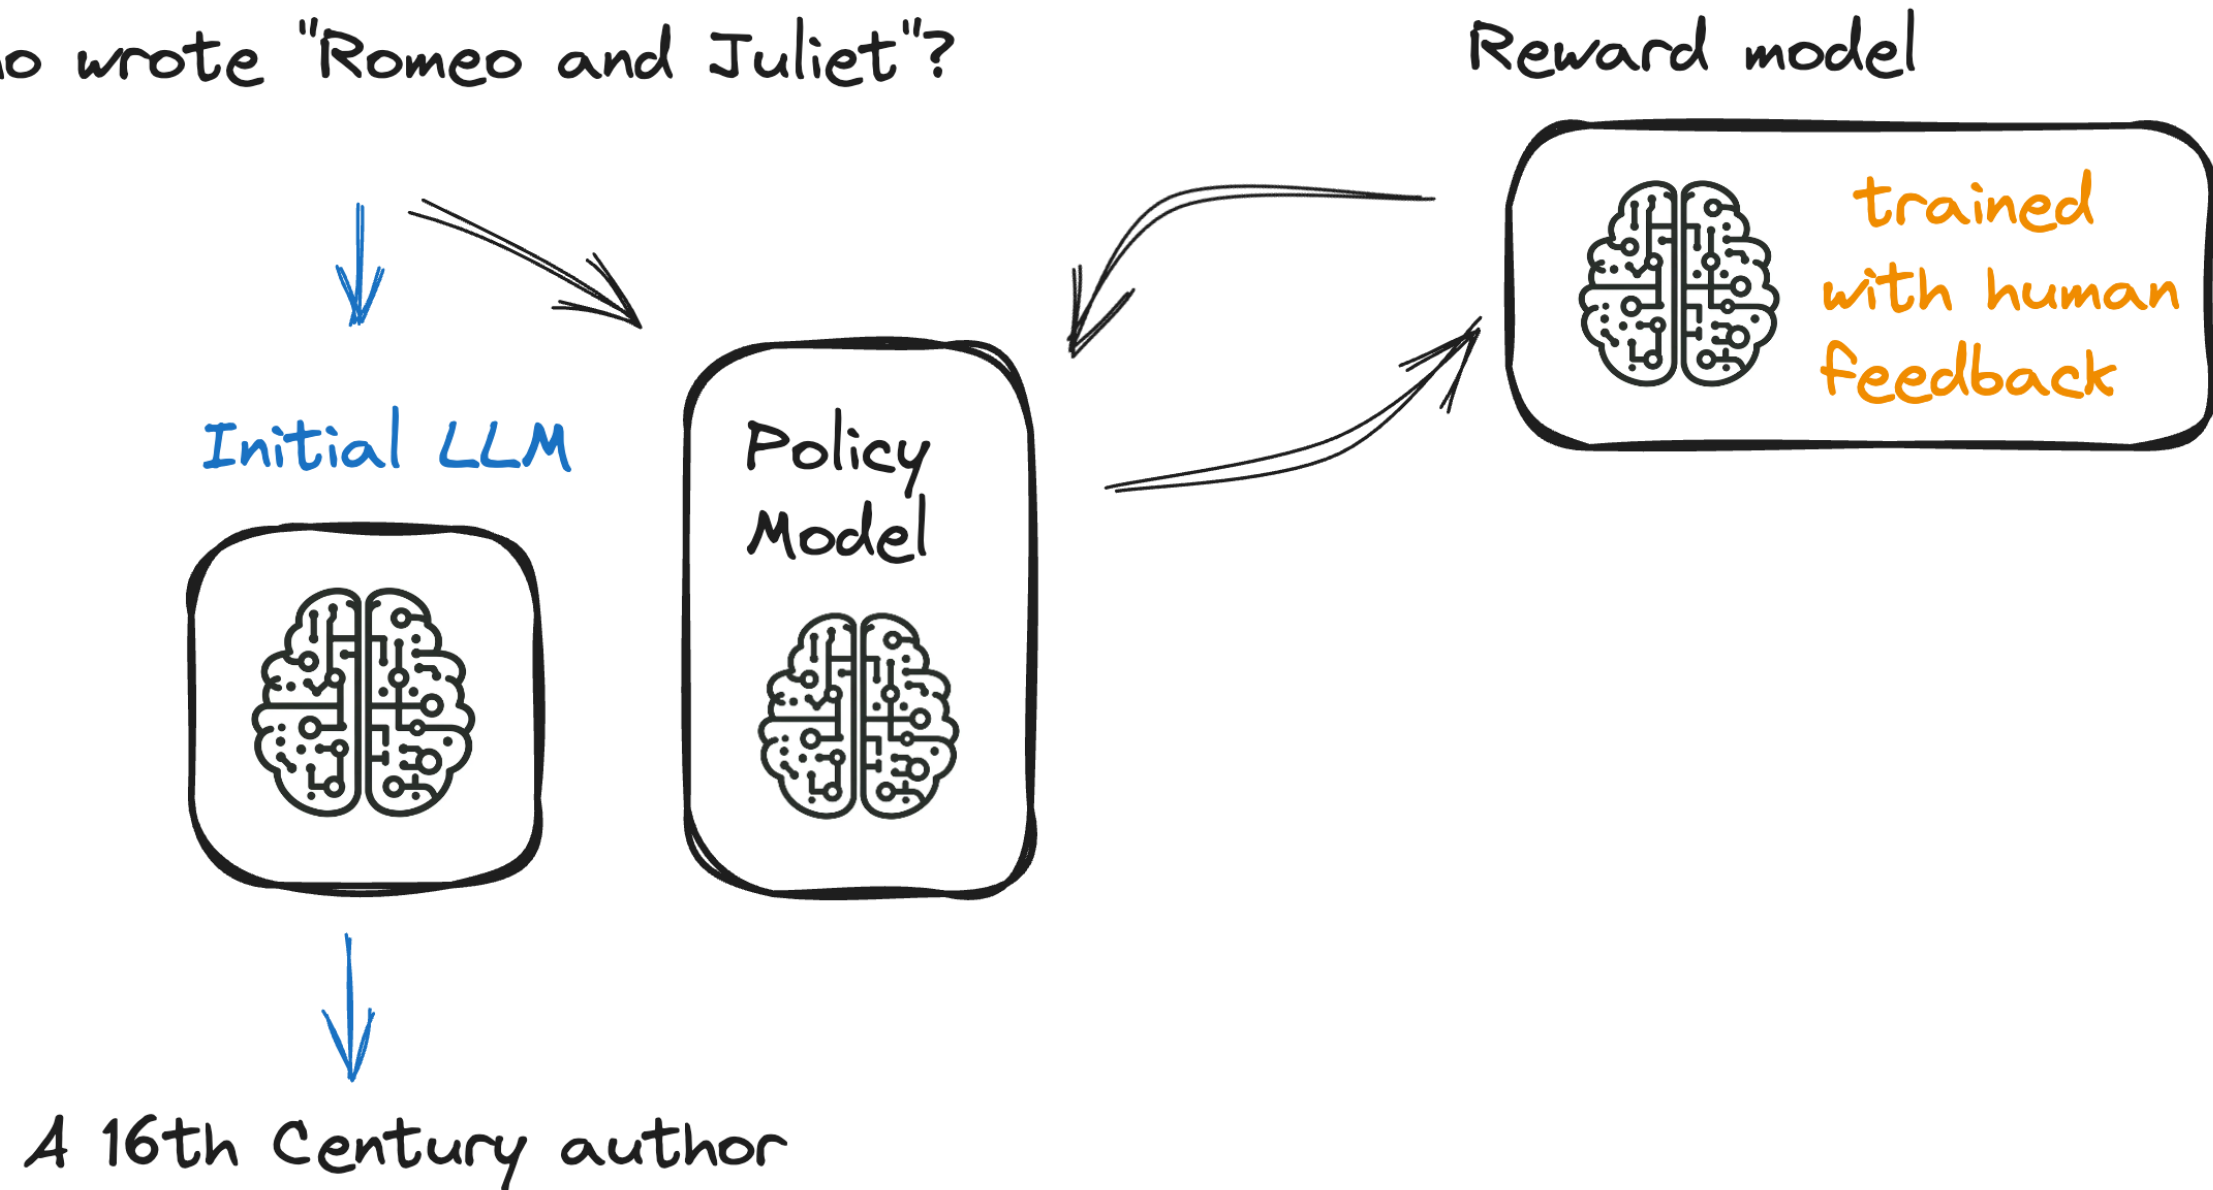
The full RLHF process

Who wrote "Romeo and Juliet"?



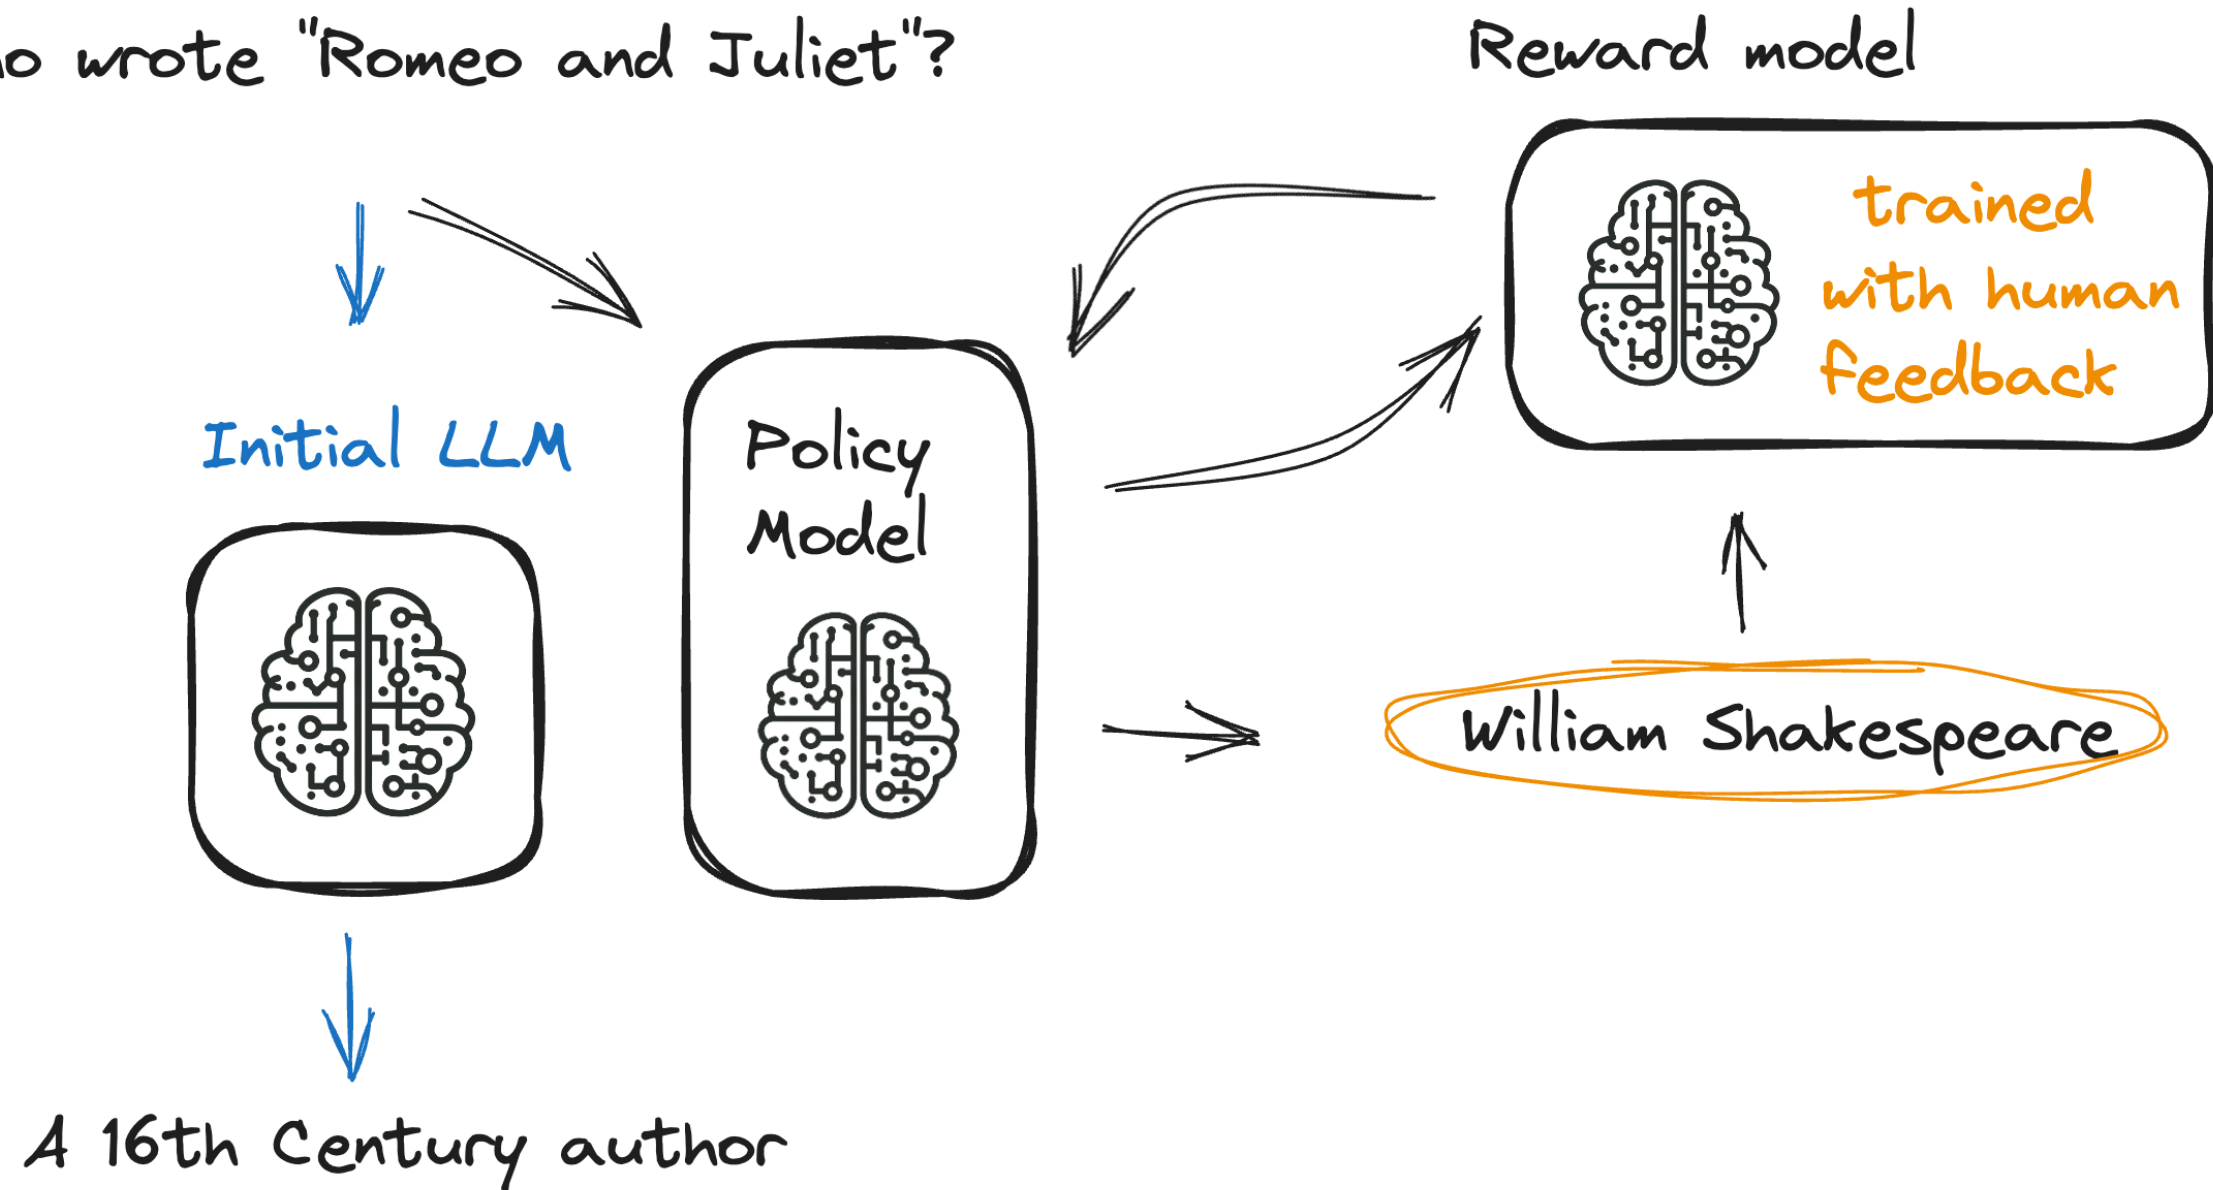
The full RLHF process

Who wrote "Romeo and Juliet"?



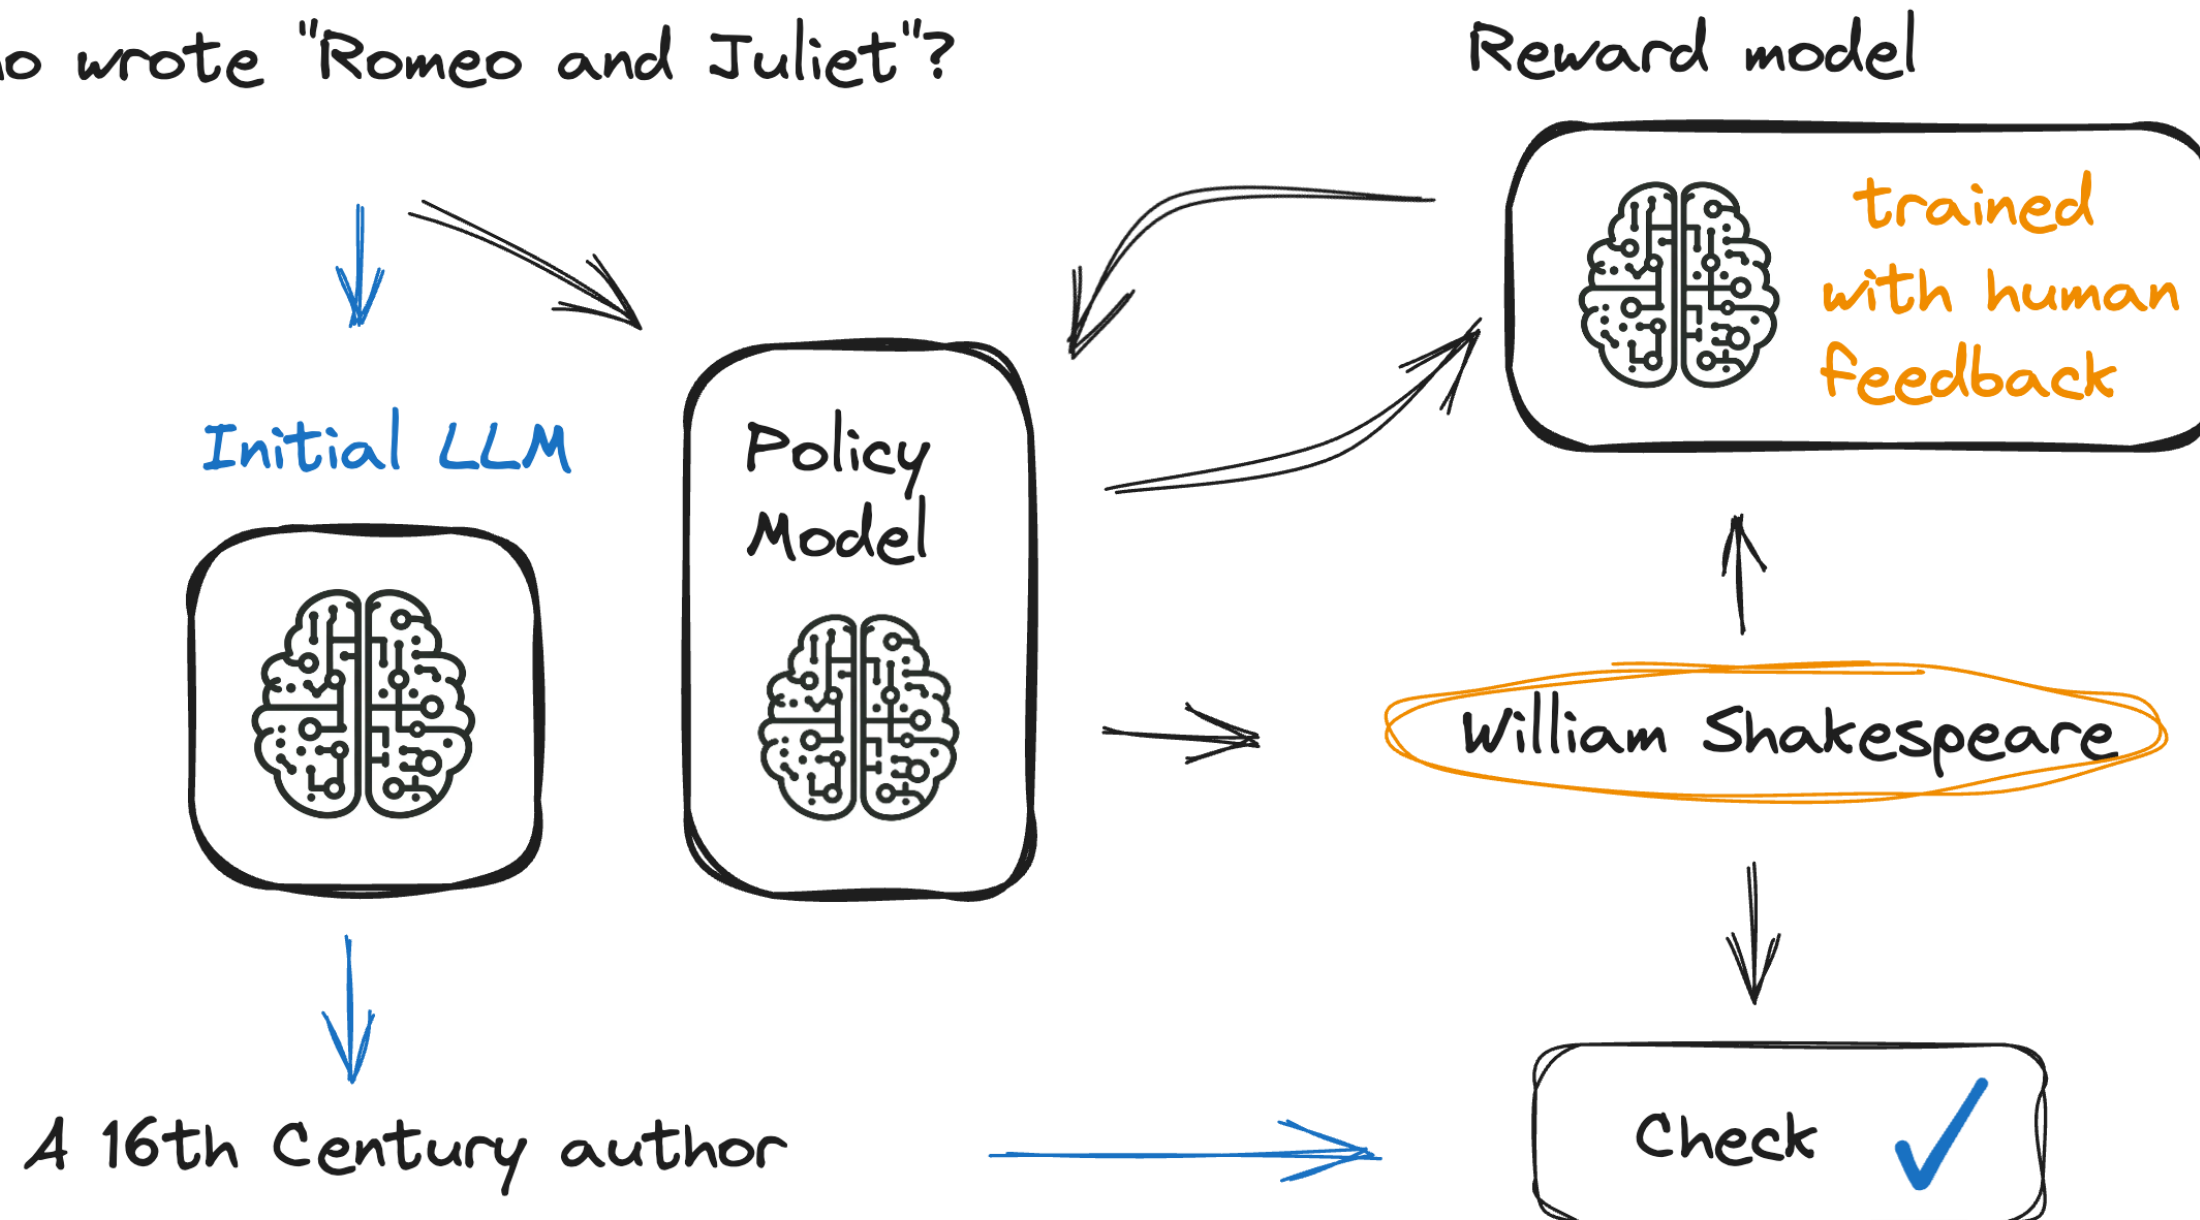
The full RLHF process

Who wrote "Romeo and Juliet"?



The full RLHF process

Who wrote "Romeo and Juliet"?



Interacting with RLHF-tuned LLMs

- Pre-trained RLHF models on Hugging Face `□`

```
from transformers import pipeline
text_generator = pipeline('text-generation', model='lvwerra/gpt2-imdb-pos-v2')
# Provide a review prompt
review_prompt = "This is definitely a"

# Generate the continuation
output = text_generator(review_prompt, max_length=50)

#Print the generated text
print(output[0]['generated_text'])
```

```
This is definitely a crucial improvement.
```

Interacting with RLHF-tuned LLMs

```
from transformers import pipeline, AutoModelForSequenceClassification, AutoTokenizer

# Instantiate the pre-trained model and tokenizer
model = AutoModelForSequenceClassification.from_pretrained("lvwerra/distilbert-imdb")
tokenizer = AutoTokenizer.from_pretrained("lvwerra/distilbert-imdb")

# Use pipeline to create the sentiment analyzer
sentiment_analyzer = pipeline('sentiment-analysis', model=model, tokenizer=tokenizer)

# Pass the text to the sentiment analyzer and print the result
sentiment = sentiment_analyzer("This is definitely a crucial improvement.")
print(f"Sentiment Analysis Result: {sentiment}")
```

positive

Let's practice!

REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)

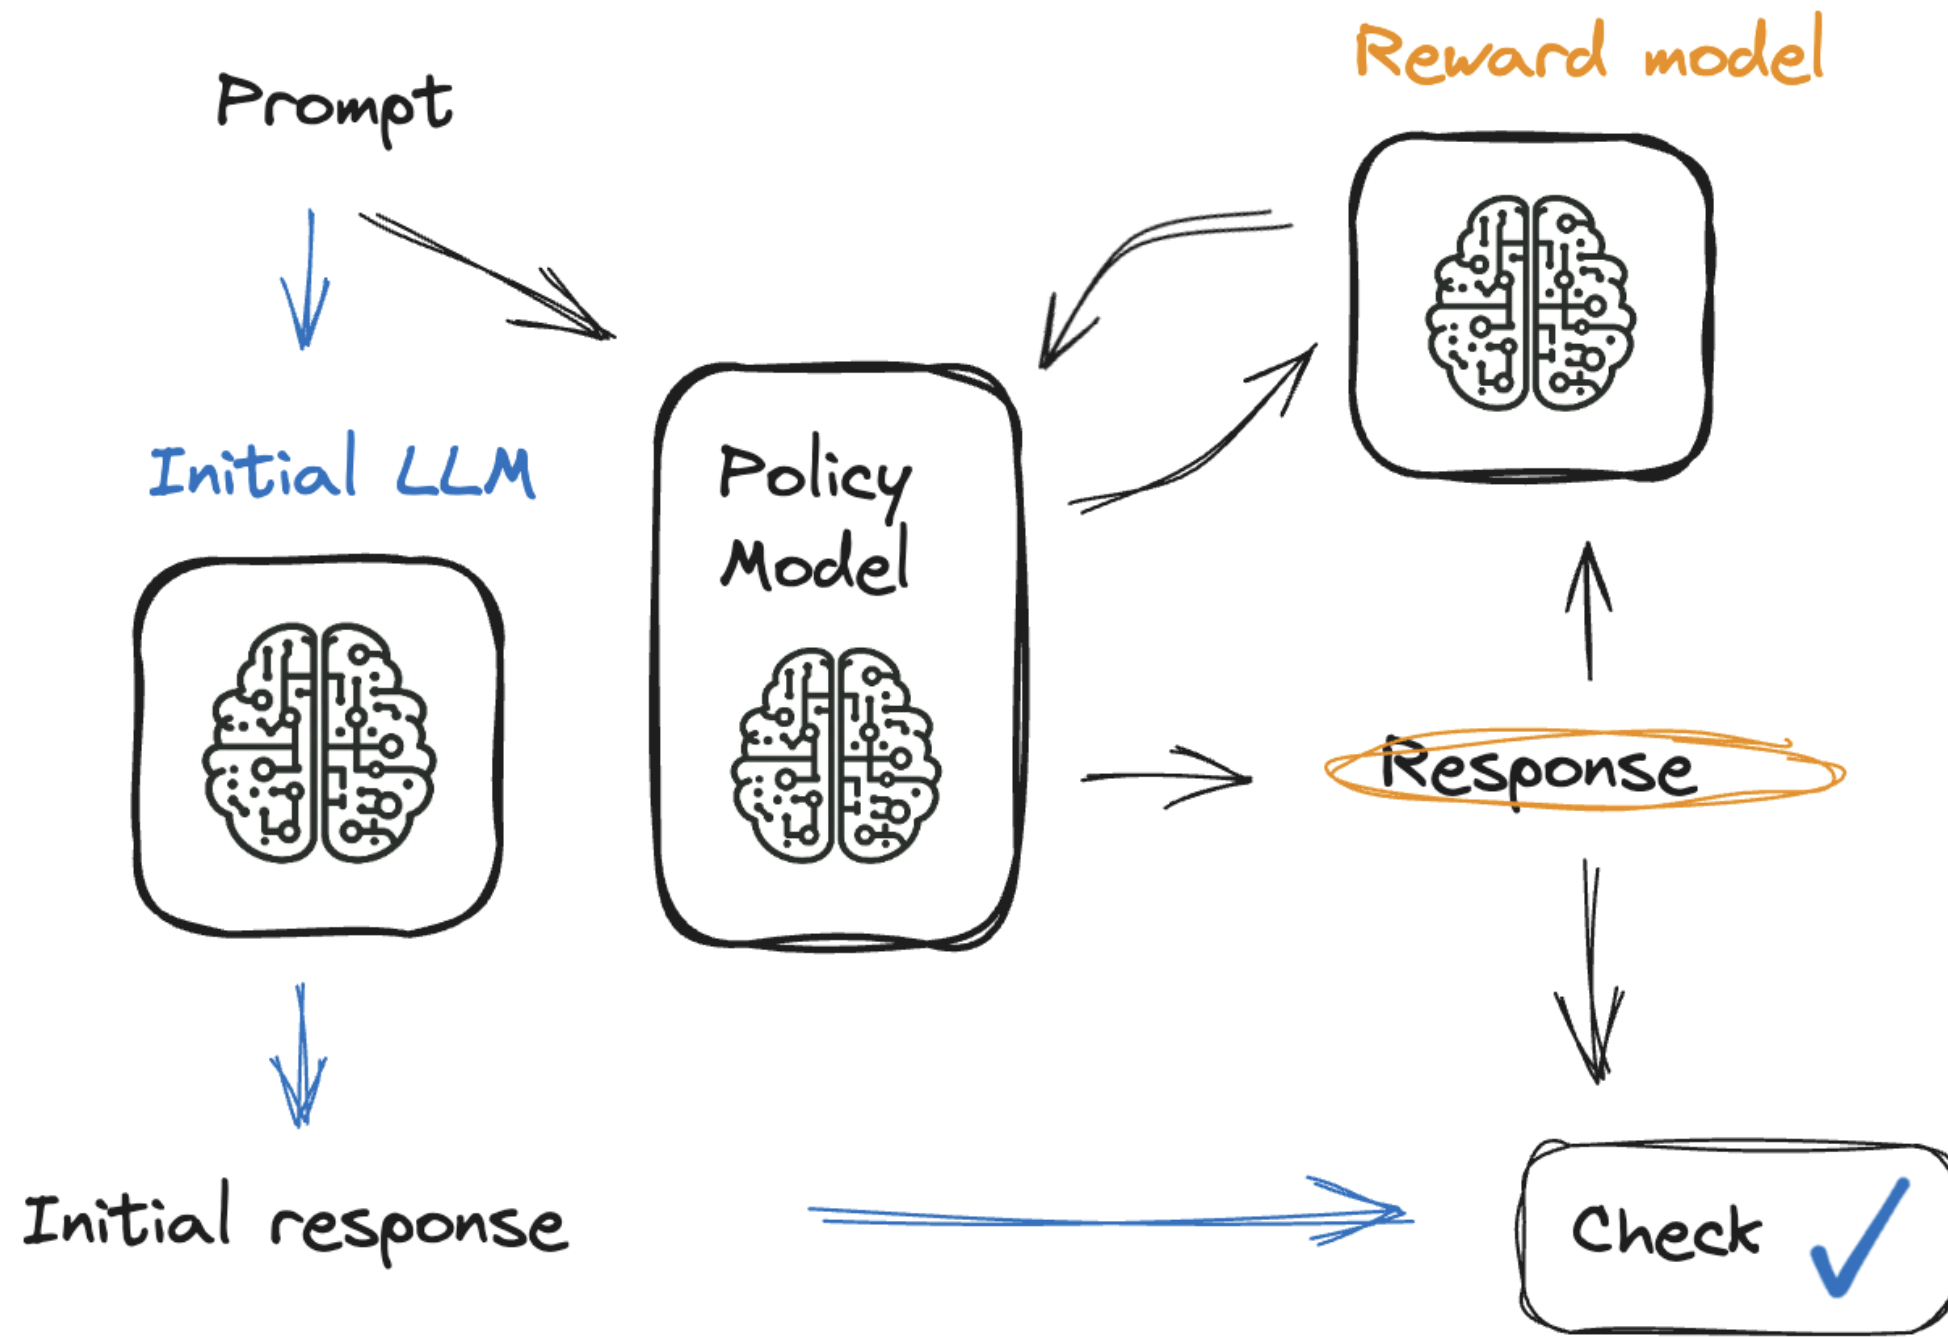
Exploring pre-trained LLMs

REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)




Mina Parham
AI Engineer

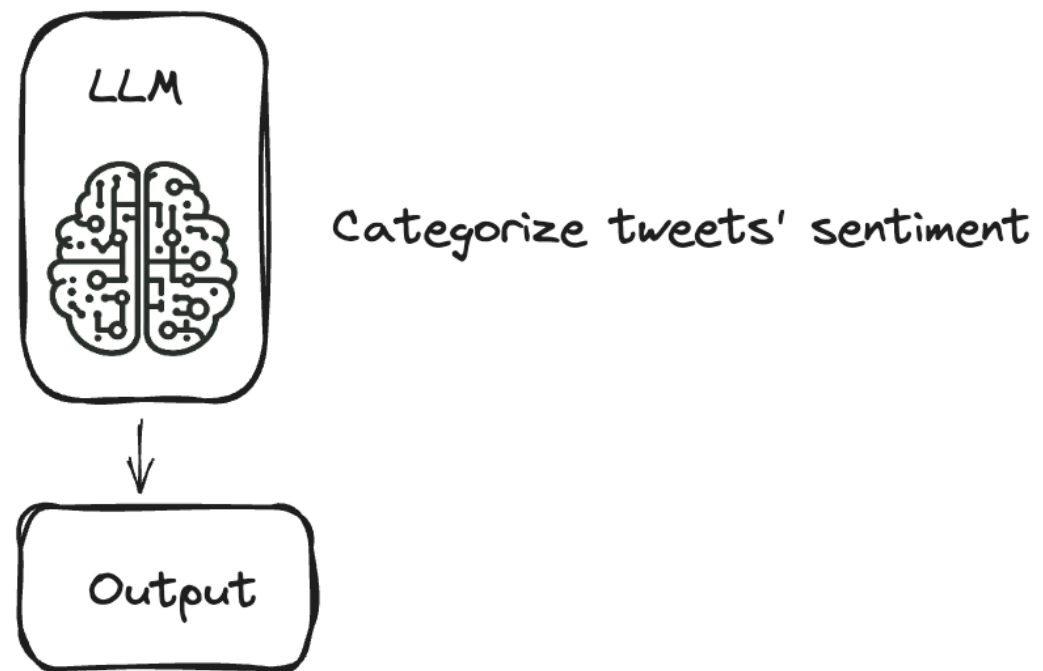
The importance of fine-tuning



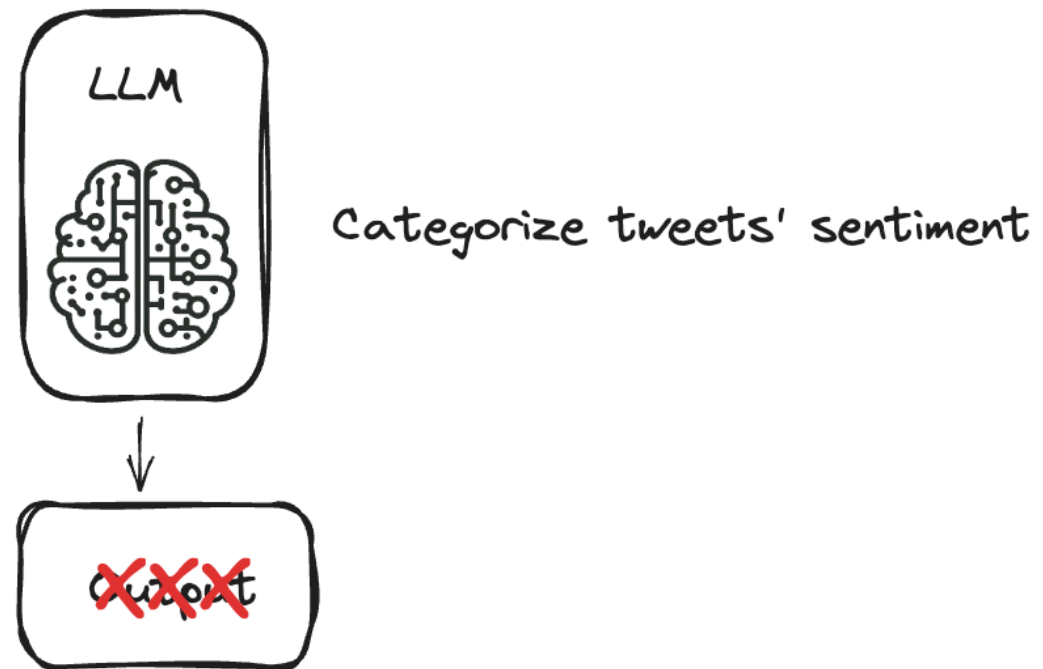
The importance of fine-tuning



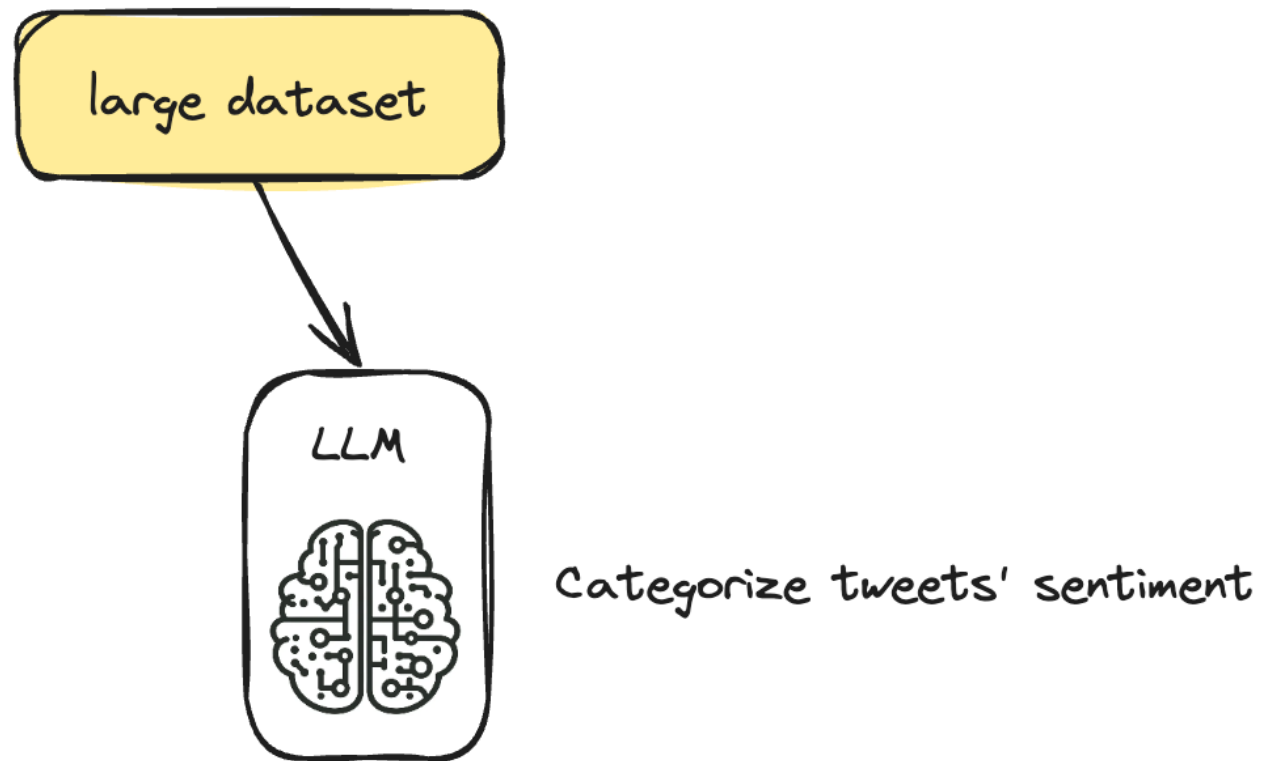
A step-by-step guide to fine-tuning an LLM



A step-by-step guide to fine-tuning an LLM



A step-by-step guide to fine-tuning an LLM



Step 1: load the data to use

```
from datasets import load_dataset
import pandas as pd

# `load_dataset` simplifies loading and preprocessing datasets from various sources
# It provides easy access to a wide range of datasets with minimal setup
dataset = load_dataset("mteb/tweet_sentiment_extraction")
df = pd.DataFrame(dataset['train'])
```

	id	text	label	label_text
0	cb774db0d1	I'd have responded, if I were going	1	neutral
1	549e992a42	Sooo SAD I will miss you in San Diego!!!	0	negative
2	08ac60f138	my boss is bullying me...	0	negative

Step 2: choose a pre-trained model

```
from transformers import AutoModelForCausalLM

# AutoModelForCausalLM simplifies loading and switching models
model = AutoModelForCausalLM.from_pretrained("openai-gpt")
```

- **Causal models:** previous tokens "cause" subsequent ones

Step 3: tokenizer

```
from transformers import AutoTokenizer

# `AutoTokenizer` loads the correct tokenizer for the specified model
tokenizer = AutoTokenizer.from_pretrained("openai-gpt")
tokenizer.add_special_tokens({'pad_token': '[PAD]'})
model.resize_token_embeddings(len(tokenizer))
```

- Padding: to have equal-sized batches of text

Step 3: tokenizer

```
def tokenize_function(examples):  
    tokenized = tokenizer(examples["content"], padding="max_length", truncation=True)  
    return tokenized  
  
tokenized_datasets = dataset.map(tokenize_function, batched=True)
```

- Batched parameter: for faster processing

Step 4: fine-tune using the Trainer method

```
training_args = TrainingArguments(  
    output_dir="test_trainer",  
    per_device_train_batch_size=1,  
    per_device_eval_batch_size=1,  
    gradient_accumulation_steps=4)
```

```
trainer = Trainer(  
    model=model,  
    args=training_args,  
    train_dataset=tokenized_dataset["train"],  
    eval_dataset=tokenized_dataset["test"])  
trainer.train()
```


Let's practice!

REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)

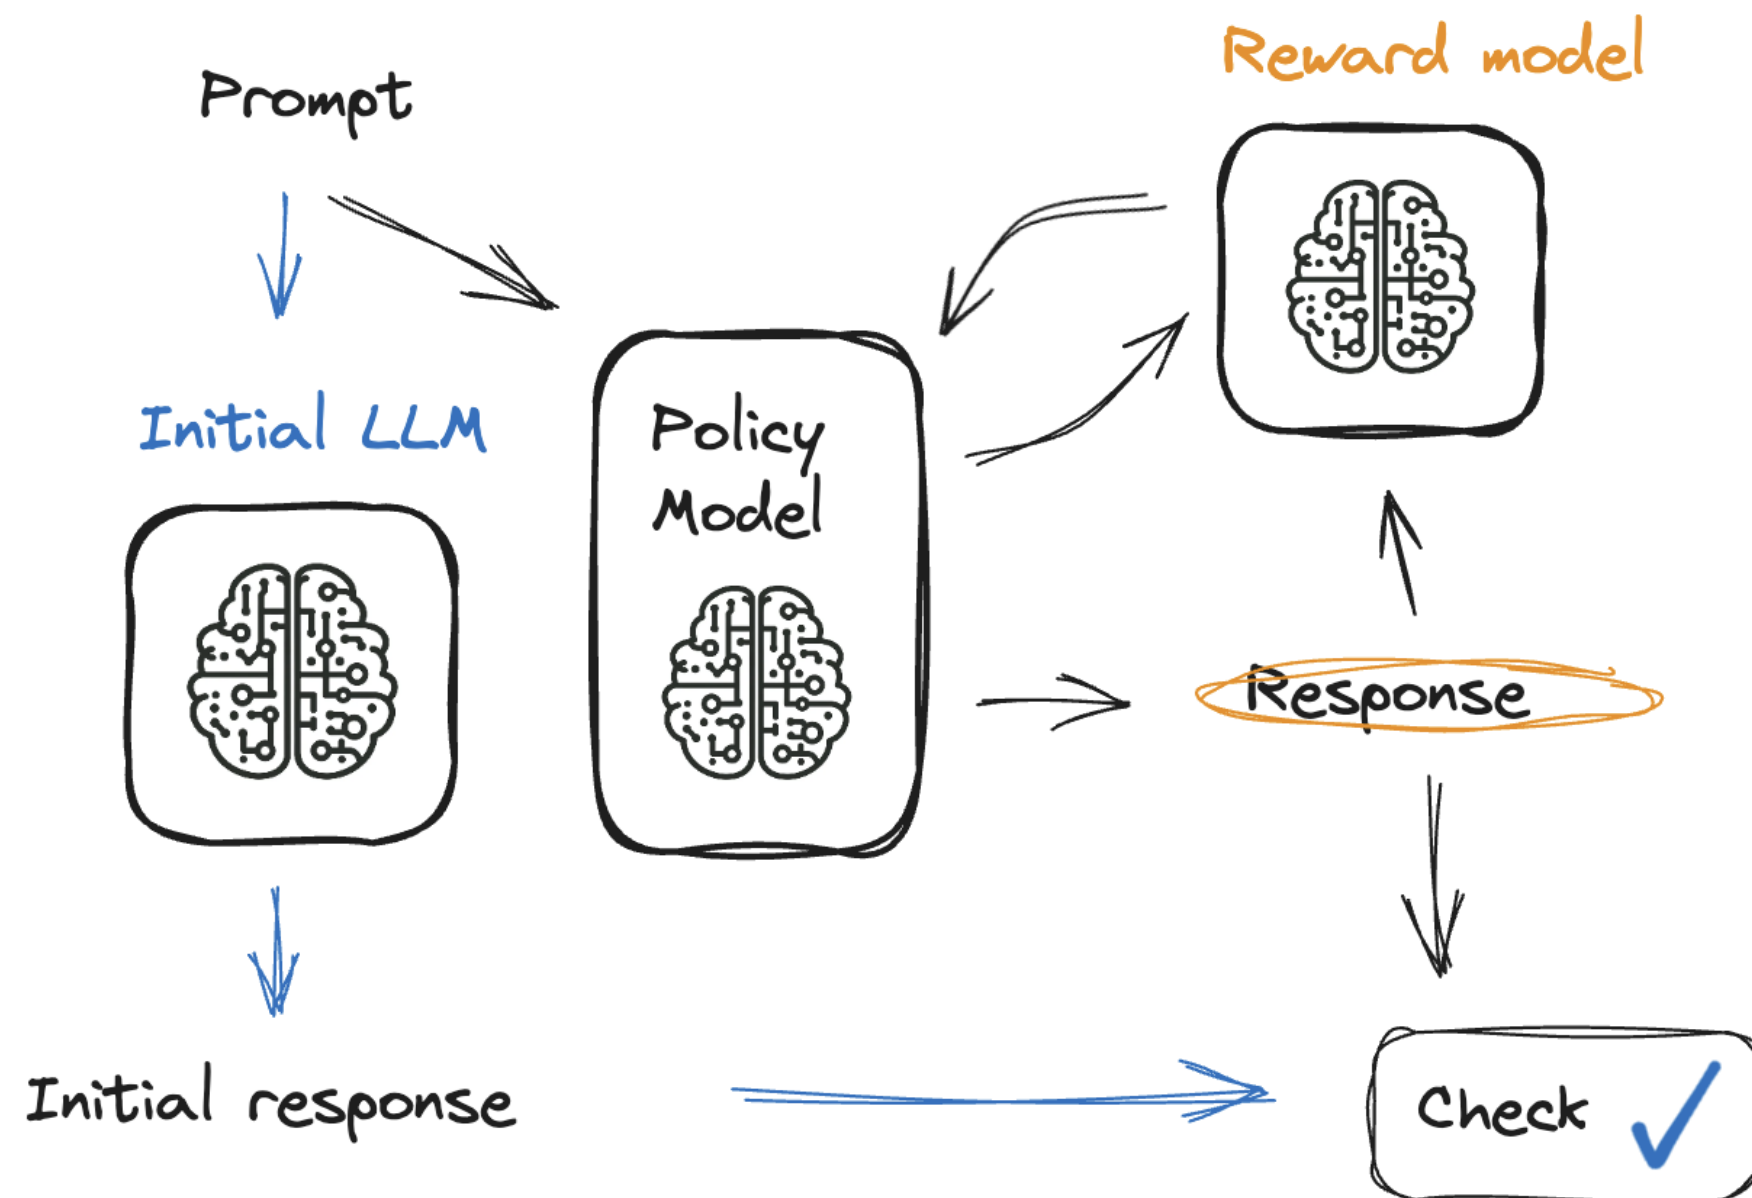
Preparing data for RLHF

REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)

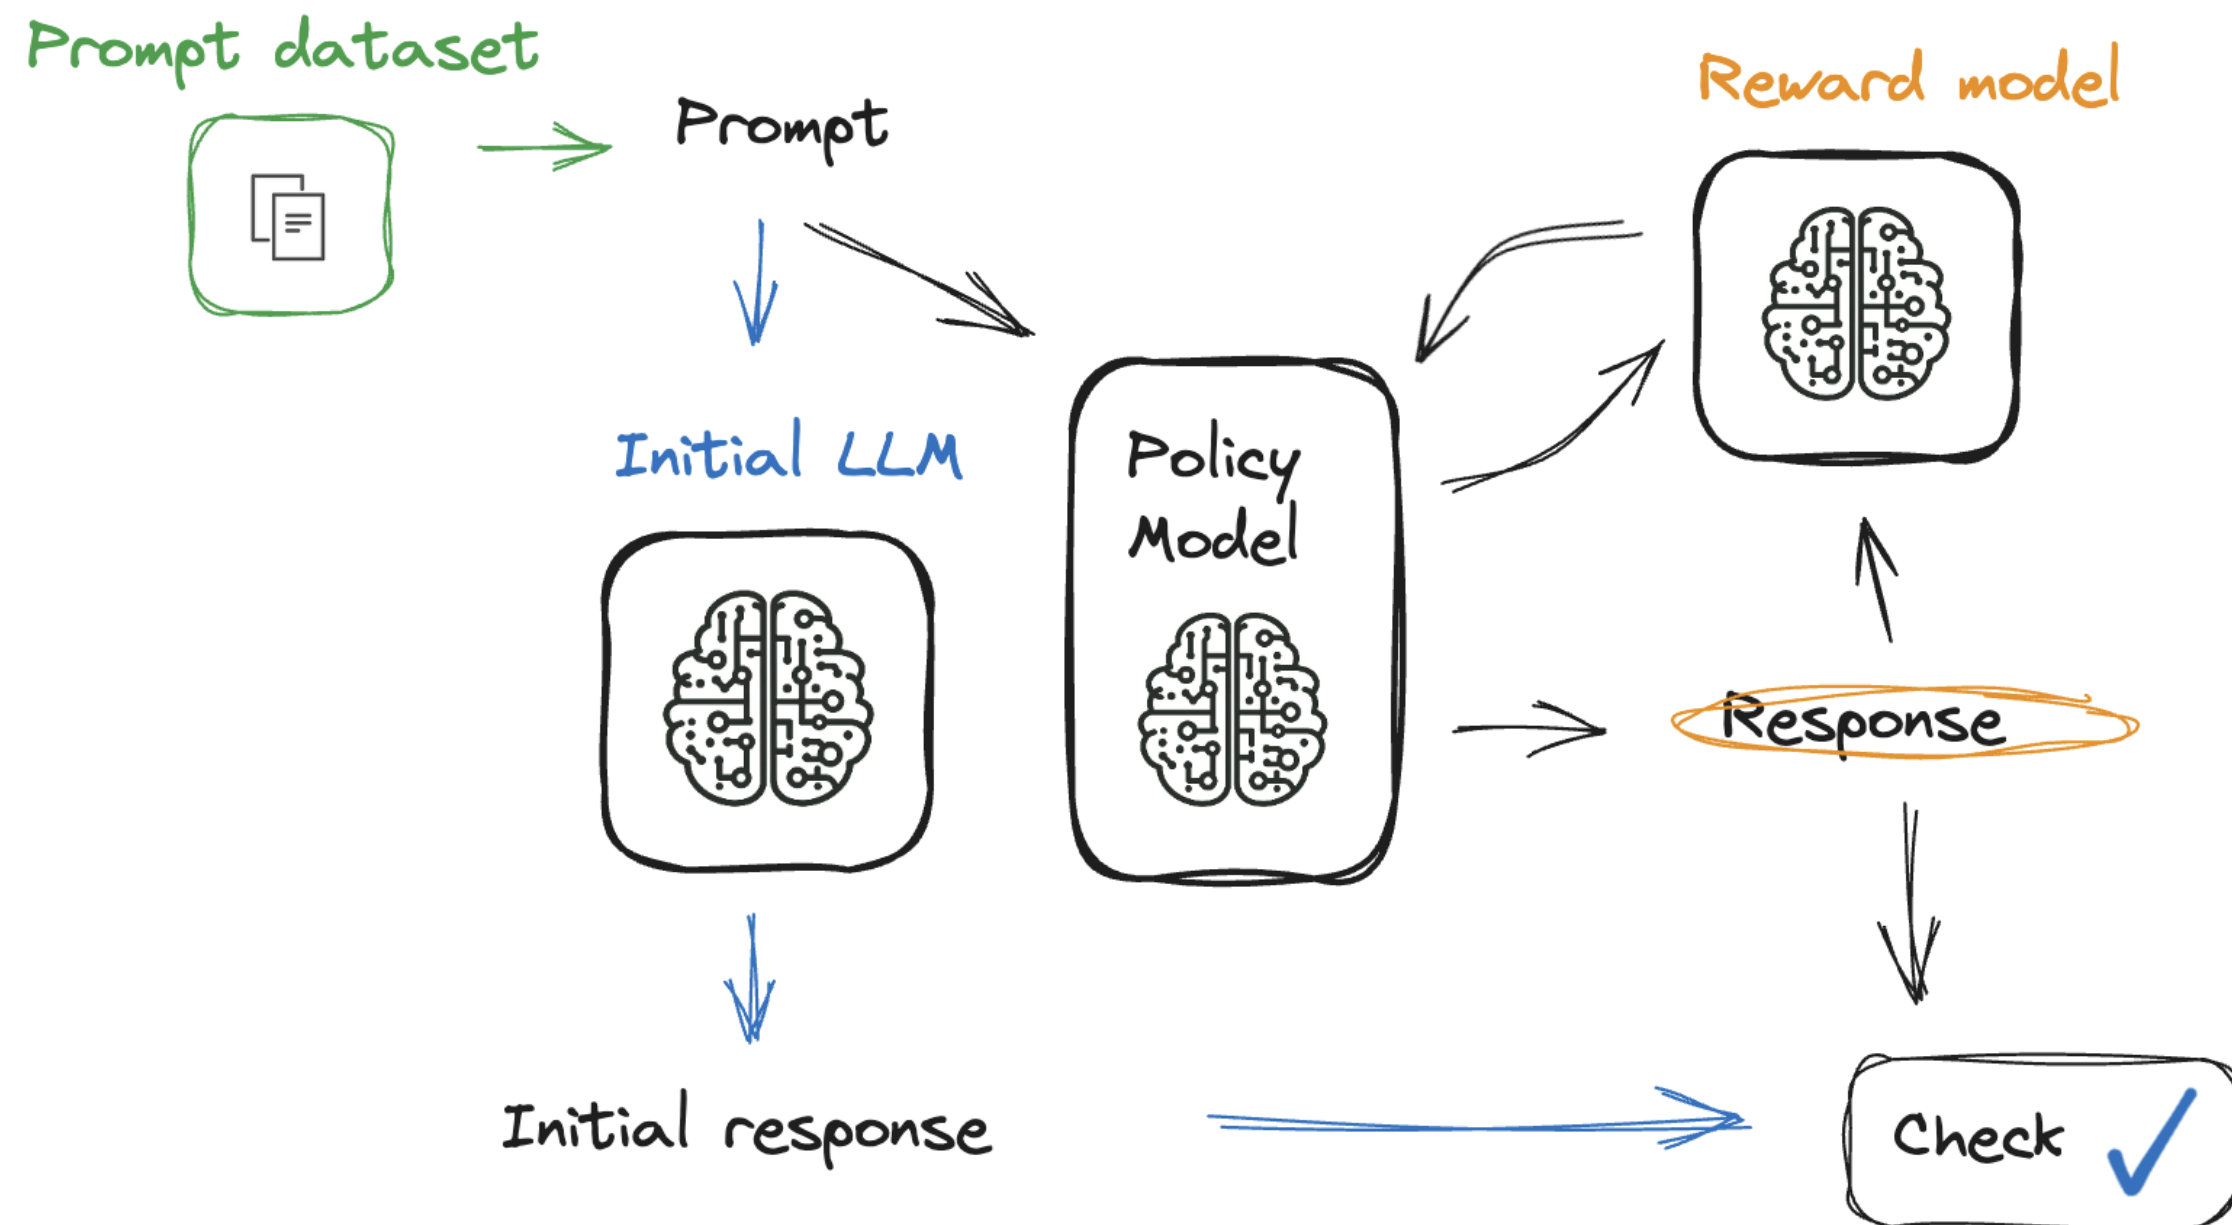


Mina Parham
AI Engineer

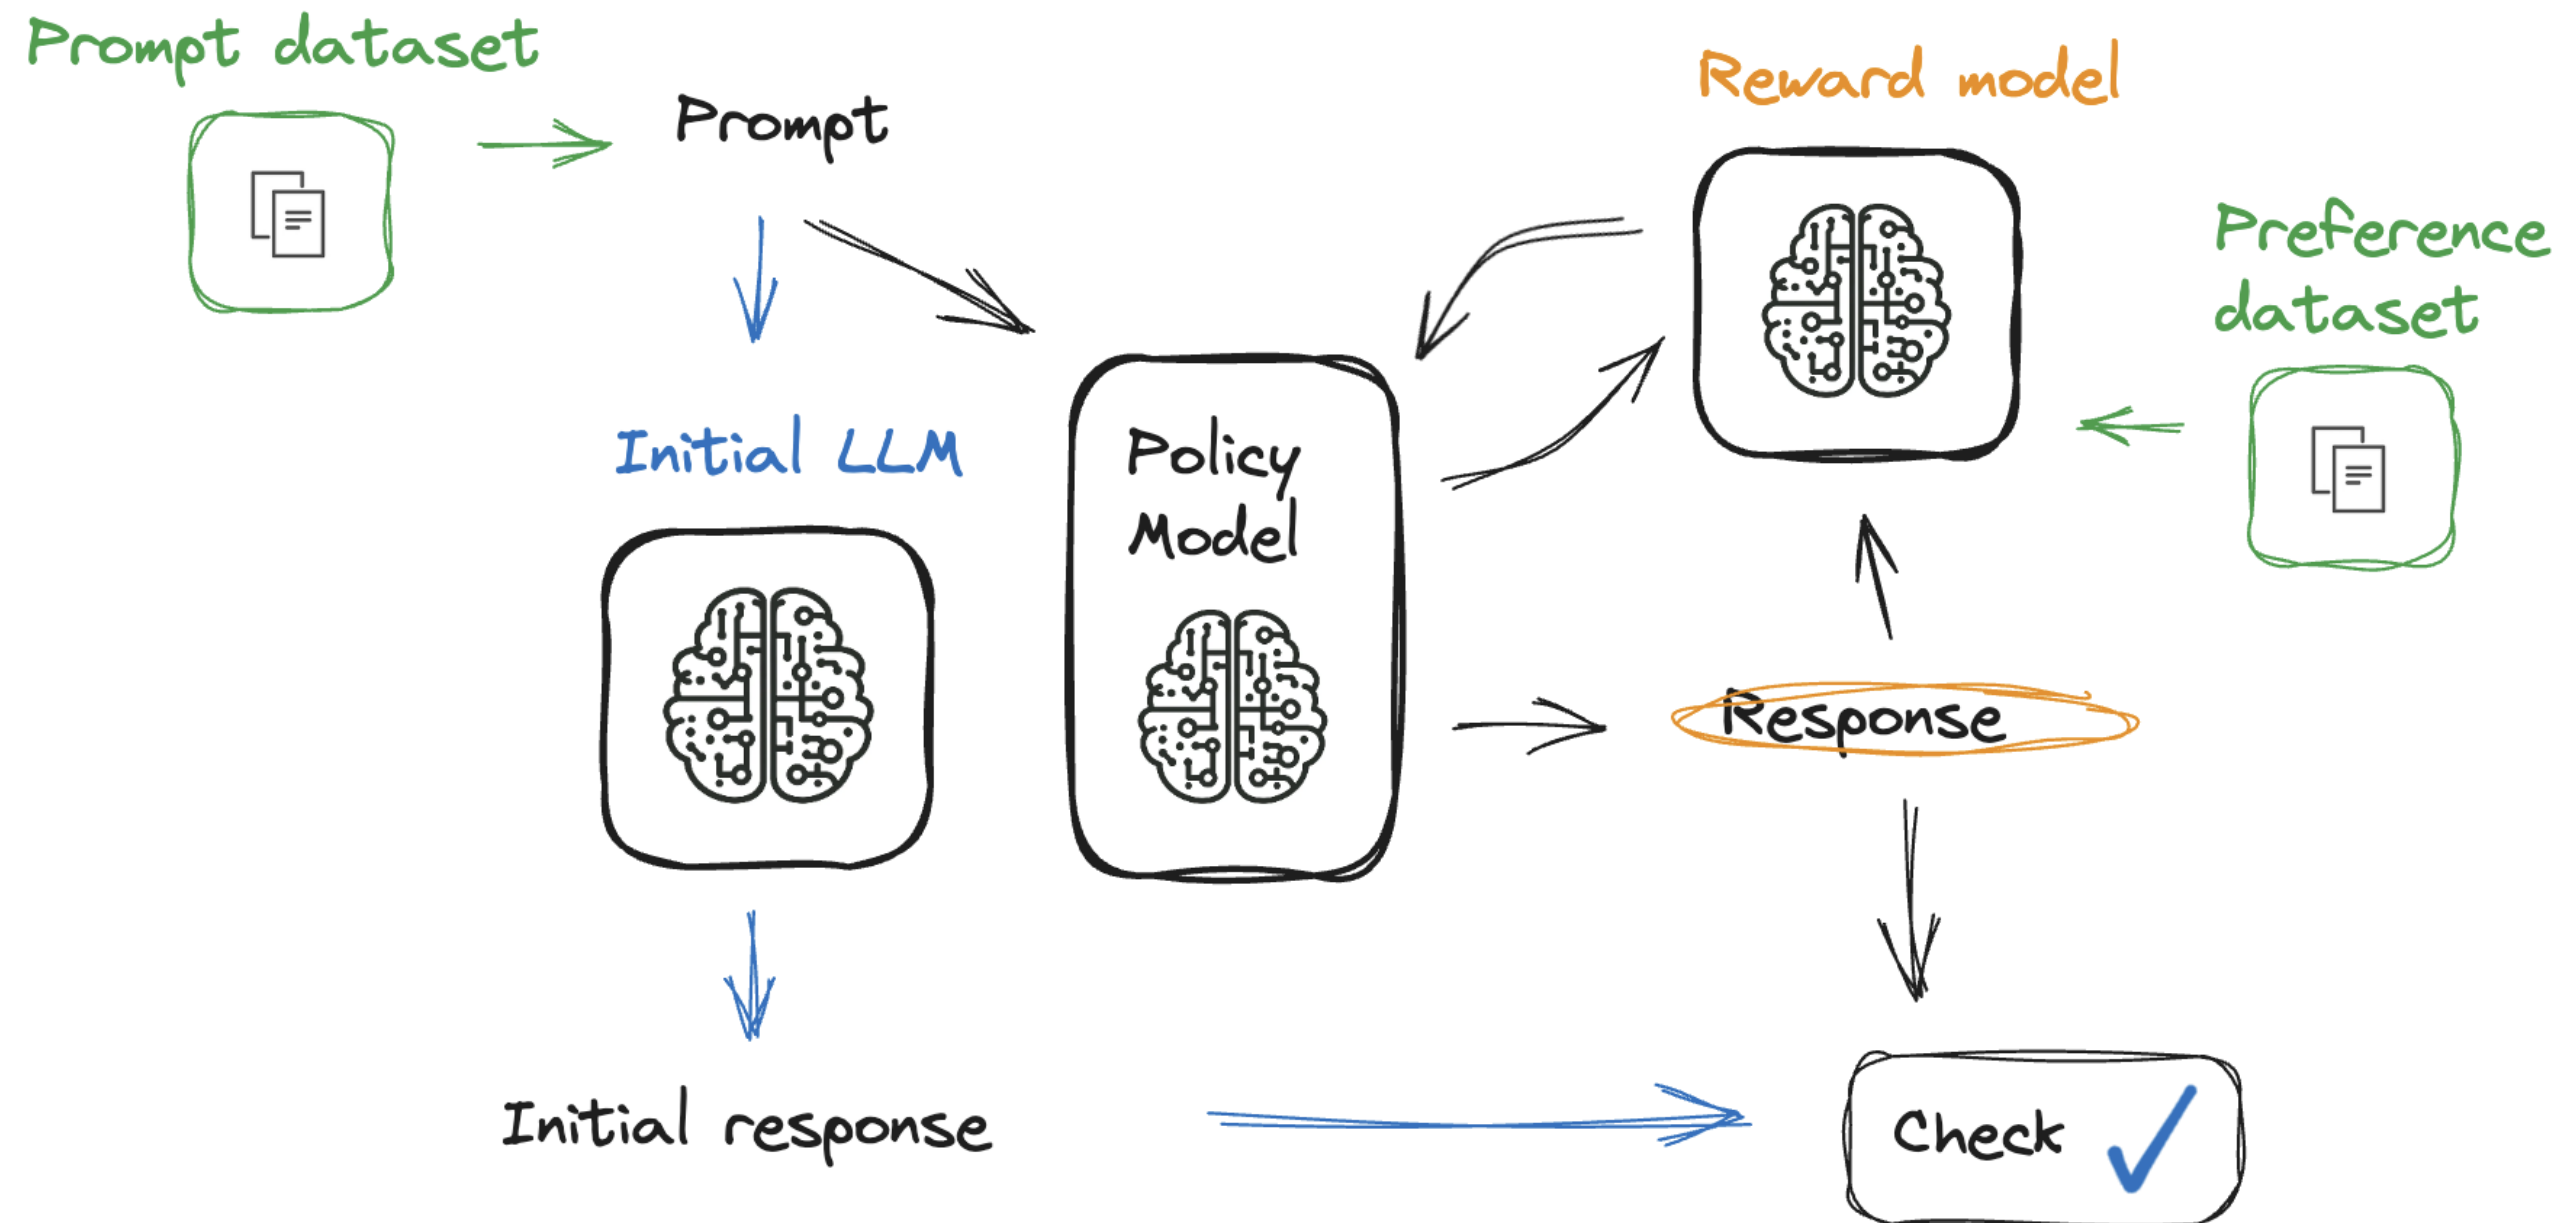
Preference vs. Prompt datasets



Preference vs. Prompt datasets



Preference vs. Prompt datasets



Prompt dataset

- Questions for the model
- Can be found on Hugging Face datasets





```
prompt_data = load_dataset("center-for-humans-and-machines/rlhf-hackathon-prompts",  
                             split="train")  
  
prompt_data['prompt'][0]
```






```
'How important is climate change?'
```





- Might need to extract the prompt
- Look for markers such as: `Input=` , `{{Text}}:` , `###Human:`






Exploring the preference dataset

```
from datasets import load_dataset
preference_data = load_dataset("trl-internal-testing/hh-rlhf-helpful-base-trl-style",
                               split="train")
```



 **Datasets:**  RLHFlow/**HH-RLHF-Helpful-standard**   like 1






Modalities:  Text Formats:  parquet Size: 100K - 1M Libraries:  Datasets  pandas  Croissant + 1

 **Dataset card**  Viewer  Files and versions  Community **1**

 **Dataset Viewer**  Auto-converted to Parquet  API  Embed  Full Screen Viewer

Split (1)
train · 115k rows

Search this dataset  SQL  Console

rejected_score  null	chosen  list · lengths	rejected  list · lengths
	 2 68	 2 68
[{"content": "What vitamins are essential for"}]		

Processing the preference dataset

```
def extract_prompt(text):  
    # Extract the prompt as the first element in the list  
    prompt = text[0]["content"]  
    return prompt
```

```
# Apply the extraction function to the dataset  
preference_data_with_prompt = preference_data.map(  
    lambda example: {**example, 'prompt': extract_prompt(example['chosen'])}  
)
```

- The way prompts are extracted is different for different datasets

Final preference dataset

```
sample = preference_data_with_prompt.select(range(1))  
sample['prompt']
```

```
'What vitamins are essential for the body to function?'
```

```
sample['chosen']
```

```
[ { "content": "What vitamins are essential for the body to function?", "role":  
  "user" }, { "content": "There are some very important vitamins that ensure the  
  proper functioning of the body, including Vitamins A, C, D, E, and K along ...} ]
```

Let's practice!

REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)