

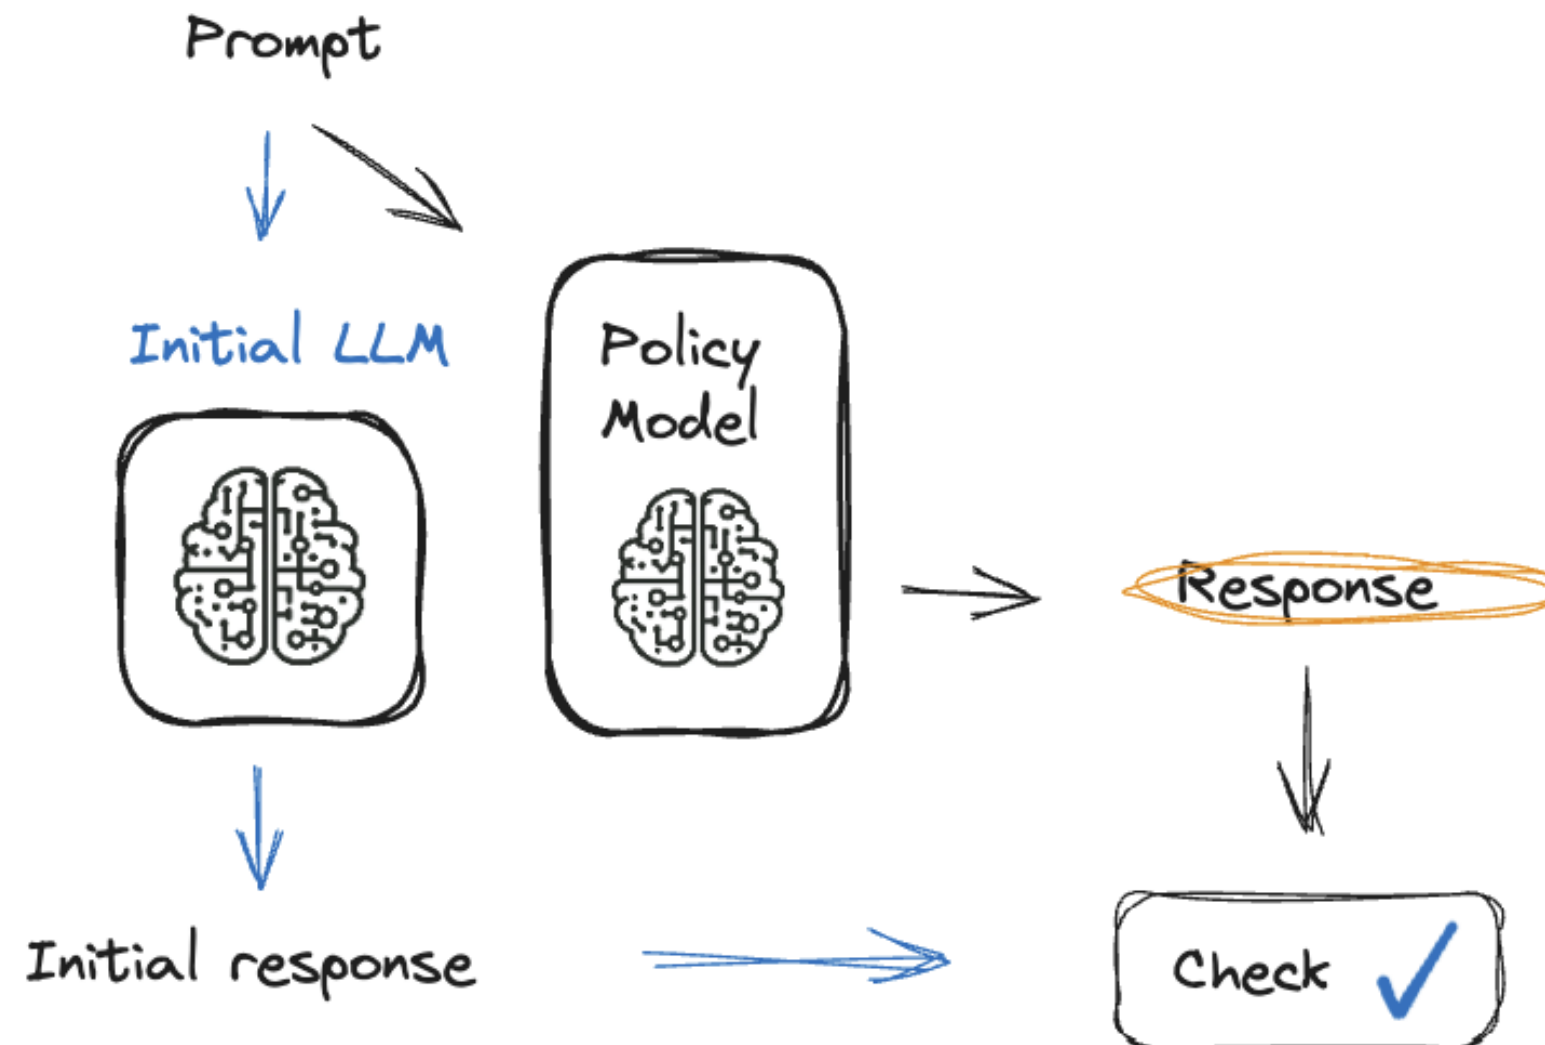
# Methods for high-quality feedback gathering

REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)

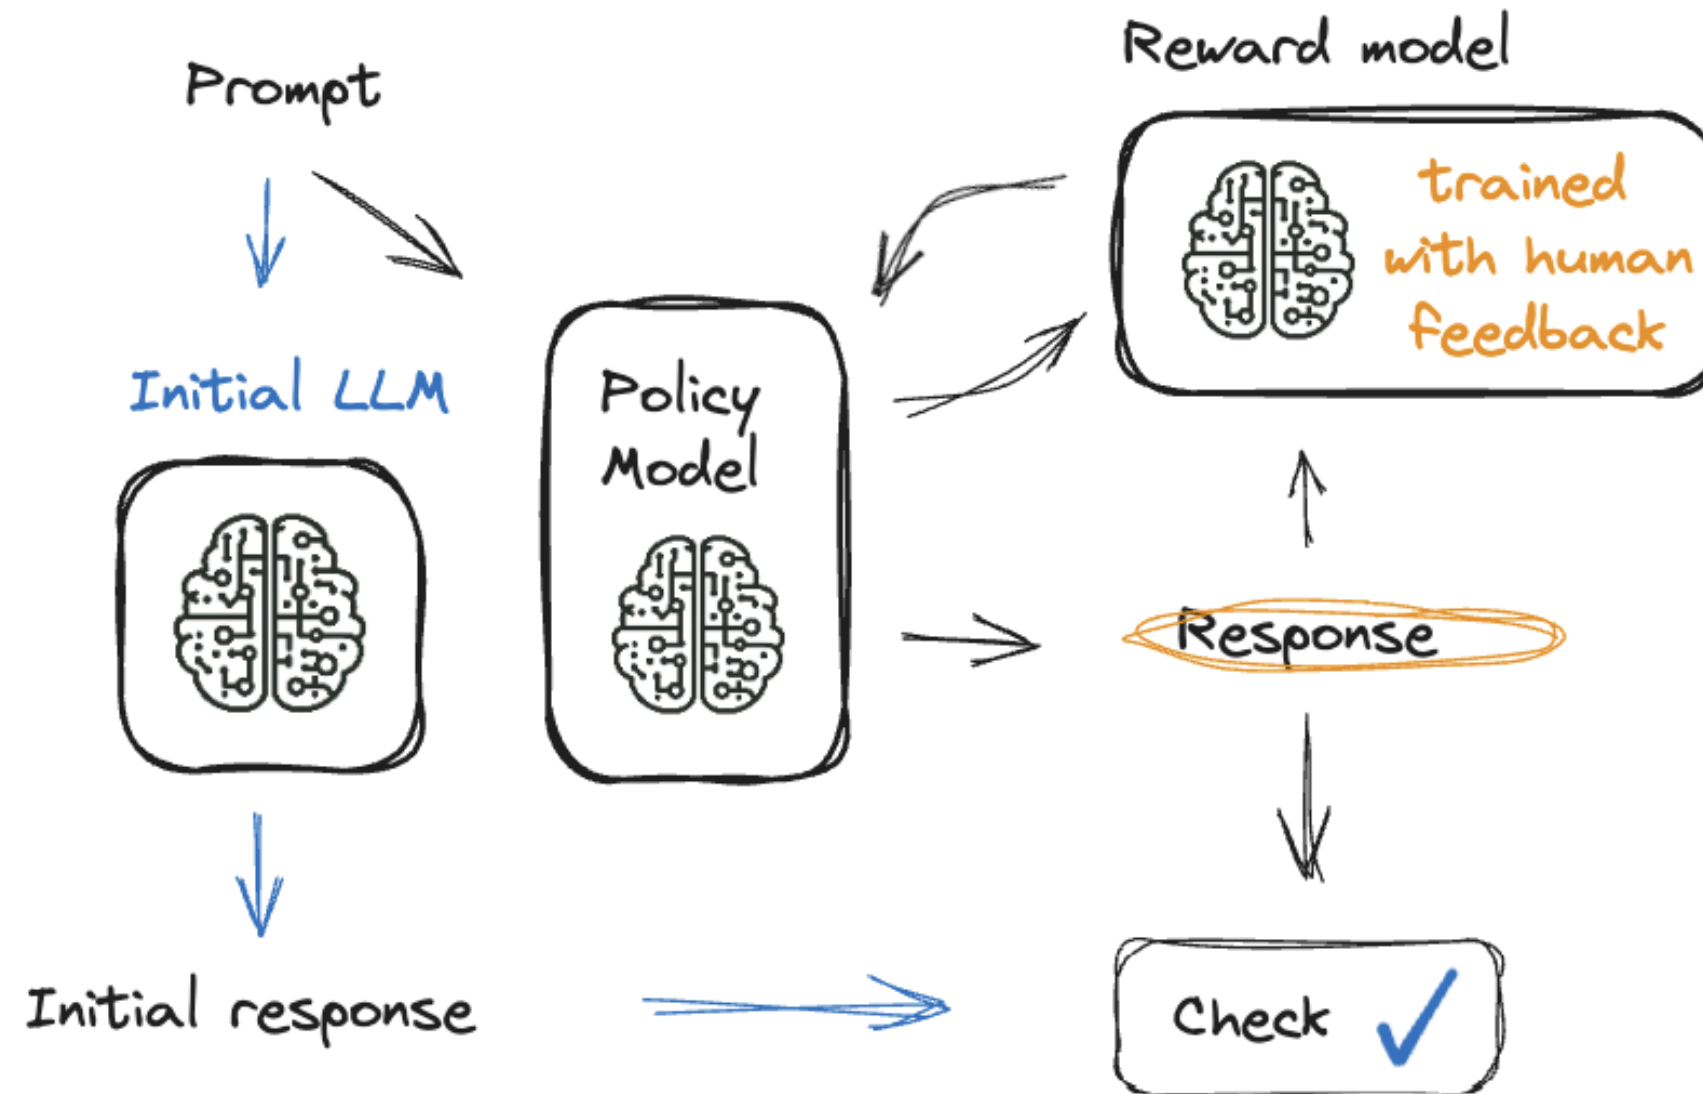


**Mina Parham**  
AI Engineer

# Methods for high-quality feedback gathering




# Methods for high-quality feedback gathering



# Pairwise comparisons

- Choosing between two options:
- **Advantages:** Simple, intuitive, reduces bias
- **Challenges:** Provides less information per label
- **Example:** Movie A vs. Movie B: "Which do you prefer?"



# Pairwise comparisons

```
def evaluate_responses(responses_A, responses_B):  
    wins_A, wins_B = 0, 0  
    for (response_A, score_A), (response_B, score_B) in zip(responses_A, responses_B):  
        if score_A > score_B:  
            wins_A += 1  
        else:  
            wins_B += 1  
    success_rate_A = (wins_A / len(responses_A)) * 100  
    success_rate_B = (wins_B / len(responses_B)) * 100  
    return success_rate_A, success_rate_B
```

# Ratings

- Assigning a score on a scale:
- **Advantages:** Provides more detailed feedback
- **Challenges:** Prone to biases, inconsistent scales
- **Example:**


Movie A: 4/5

Movie B: 3/5



# Psychological factors

- **Cognitive Biases:**
  - **Framing Effect:** How a question is presented can influence responses
  - **Serial Position Effect:** The order in which options are presented can affect decisions
  - **Anchoring:** Previous information biases current decisions





# Guidelines for collecting high-quality feedback

- Cognitive load: tired users, inconsistent feedback
- **Carefully phrase questions**
  - To combat risks from cognitive load.
- **Randomize query order**
  - To minimize bias due to anchoring and framing
- **Collect Diverse Data**
  - To mitigate the issue of noise.





# Let's practice!

REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)

# Measuring feedback quality and relevance

REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)




**Mina Parham**  
AI Engineer

# Application of detecting anomalous feedback

For example:

- **Positive Review:**
  - "I loved this product!"
- **Negative Review:**
  - "Awful service."
- **Neutral Review:**
  - "Does what it's supposed to."
- **Outlier Review:**
  - "The sky is blue."



# Detecting anomalous feedback

```
import numpy as np
```

```
def least_confidence(prob_dist):  
    simple_least_conf = np.nanmax(prob_dist)  
    num_labels = float(prob_dist.size) # number of labels  
    least_conf = (1 - simple_least_conf) * (num_labels / (num_labels - 1))  
    return least_conf
```

```
def filter_low_confidence_predictions(prob_dists, threshold=0.5):  
    filtered_indices = [i for i, prob_dist in enumerate(prob_dists)  
                        if least_confidence(prob_dist) > threshold]  
    return filtered_indices
```

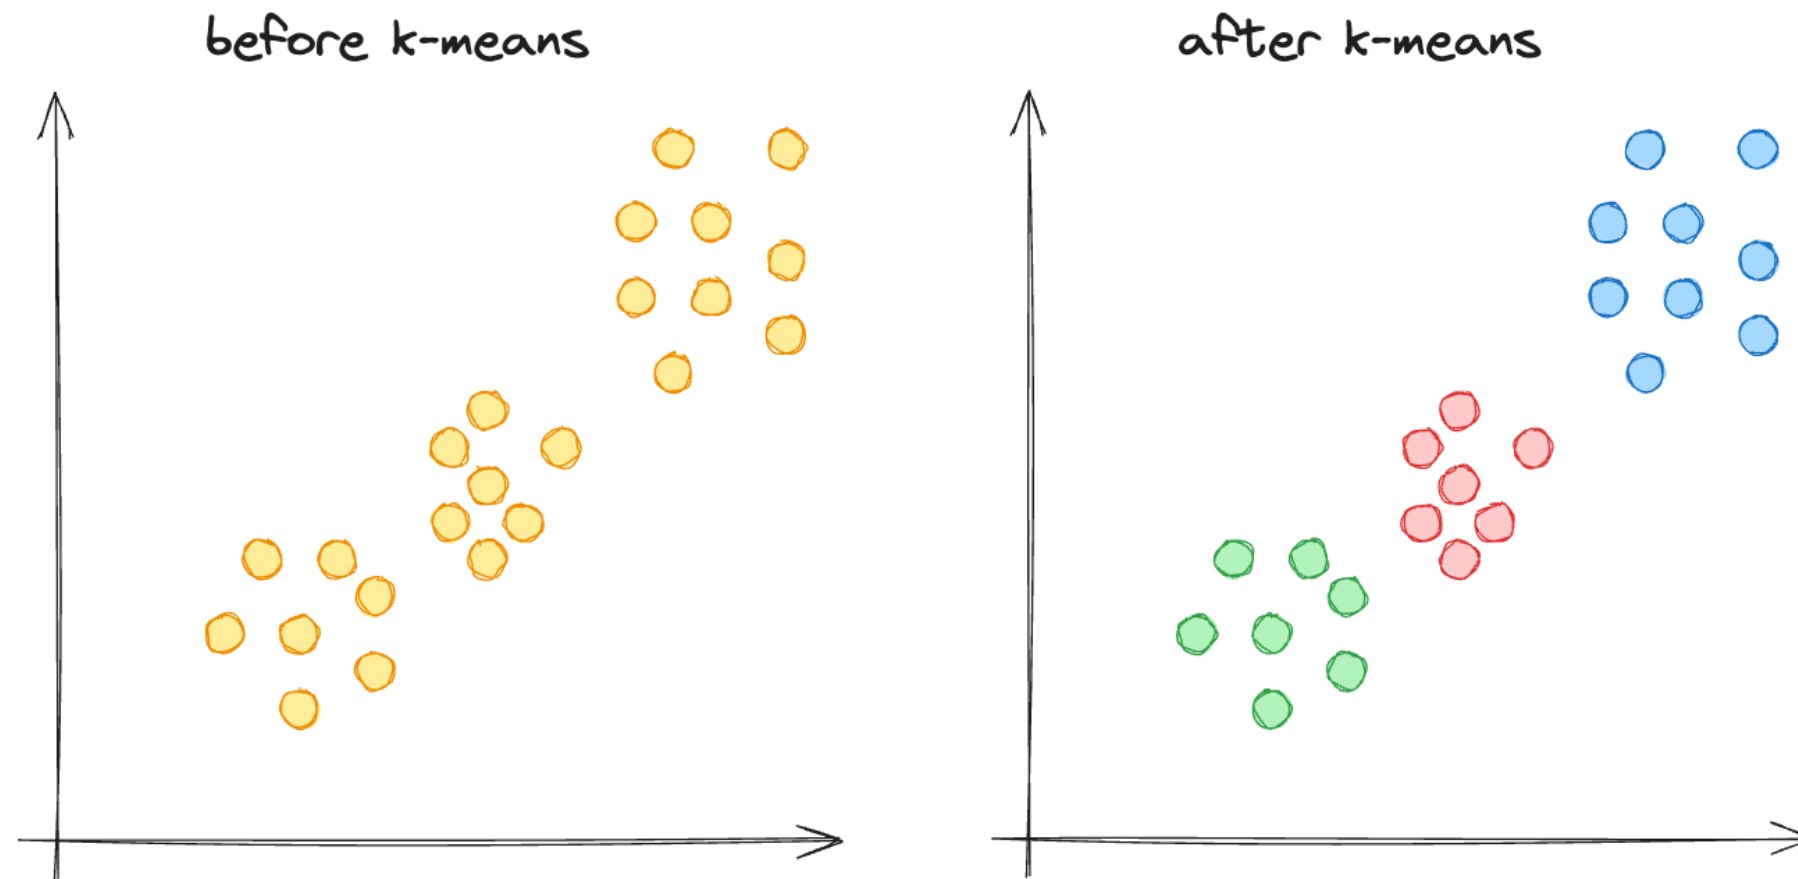
# Detecting anomalous feedback

```
prob_distribution_array = np.array([
    [0.1, 0.1, 0.2],    # Low confidence (0.2)
    [0.6, 0.2, 0.1],    # High confidence (0.6)
    [0.3, 0.3, 0.4]     # Medium confidence (0.4)
])
# Filter function with 0.5 threshold
filtered_feedback_indices, filtered_confidences =
filter_low_confidence_predictions(prob_distribution_array, threshold=0.5)
print(f"Filtered Confidence Scores: {filtered_confidences}")
```

```
Filtered Confidence Scores: [0.6]
```

# K-means

- Great for detecting anomalies and quick to implement
- Use domain knowledge or analytical methods to determine number of clusters





# Anomaly detection with k-means

```
import numpy as np
import pandas as pd
from sklearn.cluster import KMeans

def detect_anomalies(data, n_clusters=3):
    kmeans = KMeans(n_clusters=n_clusters, random_state=42)
    clusters = kmeans.fit_predict(data)
    centers = kmeans.cluster_centers_
    # Calculate distances from cluster centers
    distances = np.linalg.norm(data - centers[clusters], axis=1)
    return distances
```

# Anomaly detection with k-means

```
feedback_data = np.array([
    [4.0], # Close to center of cluster
    [4.5], # Close to center of cluster
    [1.0], # Anomaly - far from main group
    [4.1], # Close to center of cluster
    [3.9]  # Close to center of cluster
])
anomalies = detect_anomalies(confidences, n_clusters=1)
print(anomalies)
```


```
[0.5 1.  2.5  0.6 0.4]
```

# Let's practice!

REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)


# Active learning

REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)




**Mina Parham**  
AI Engineer

# Human in the loop systems




# Human in the loop systems

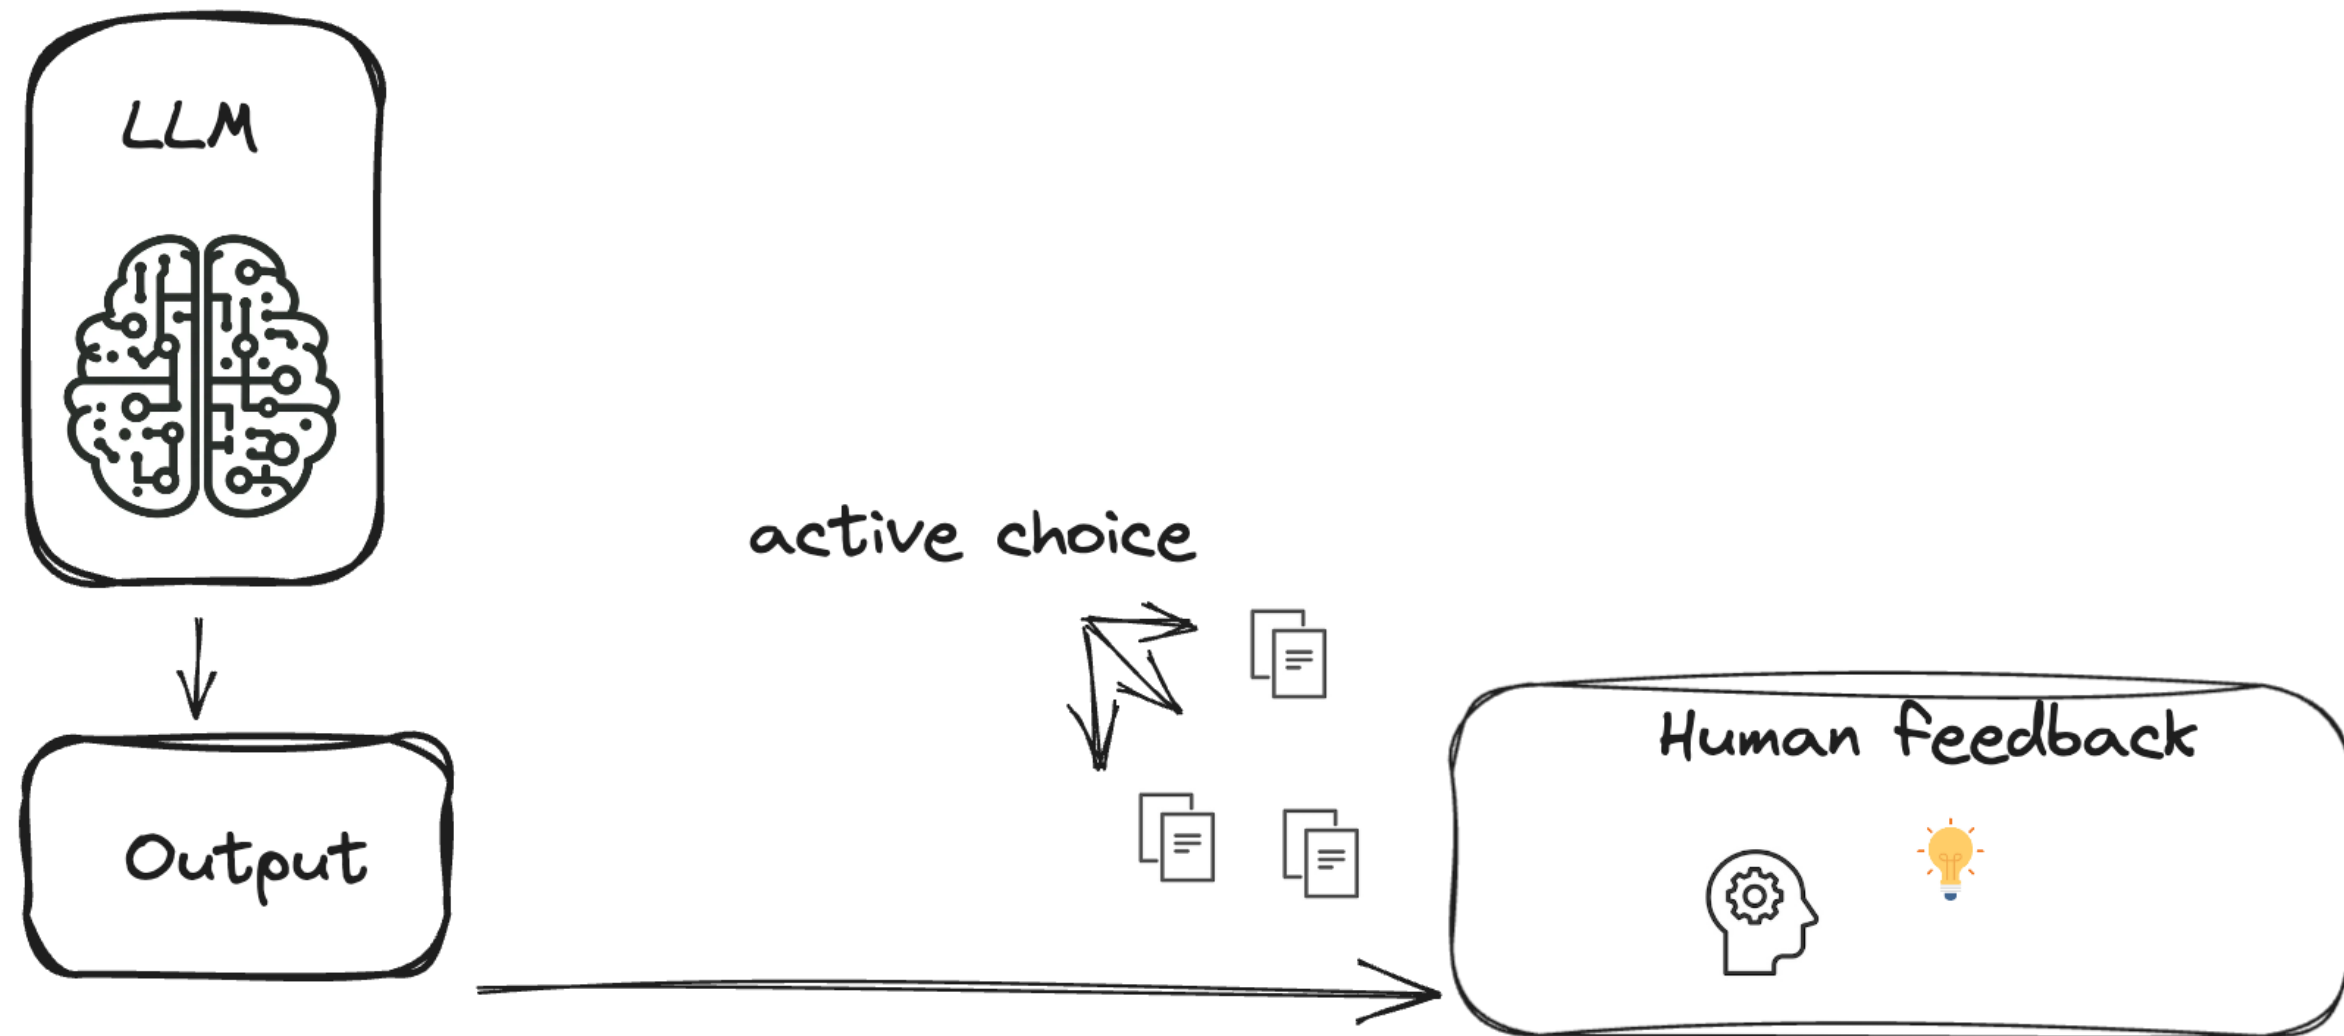




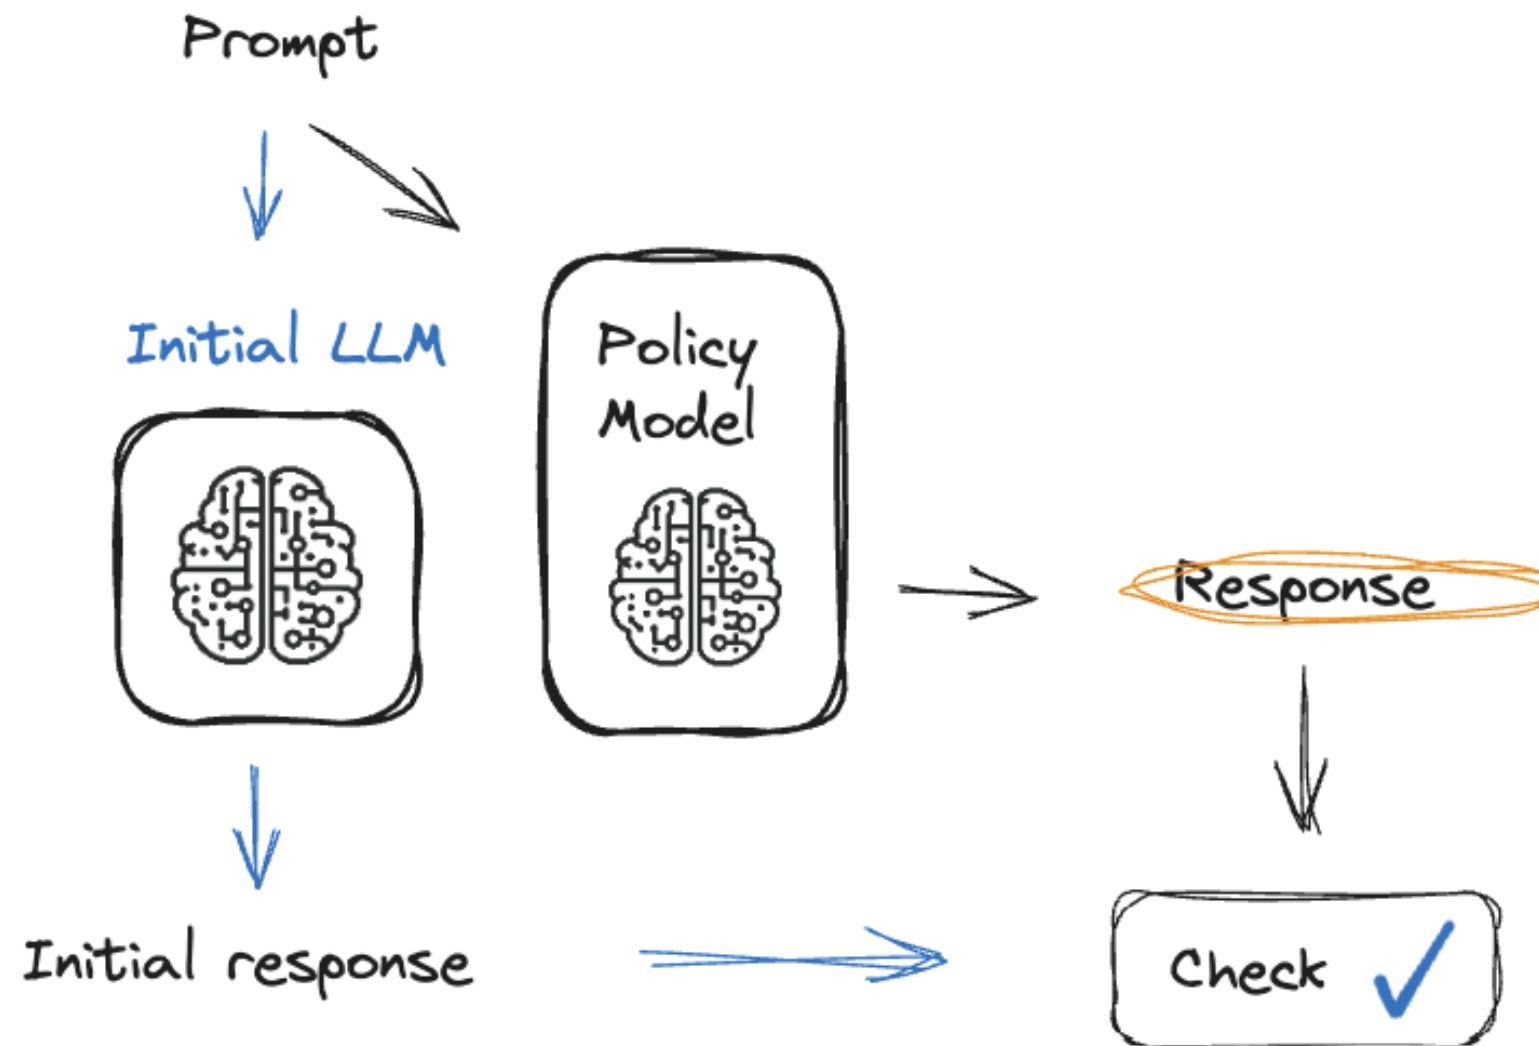
# Human in the loop systems



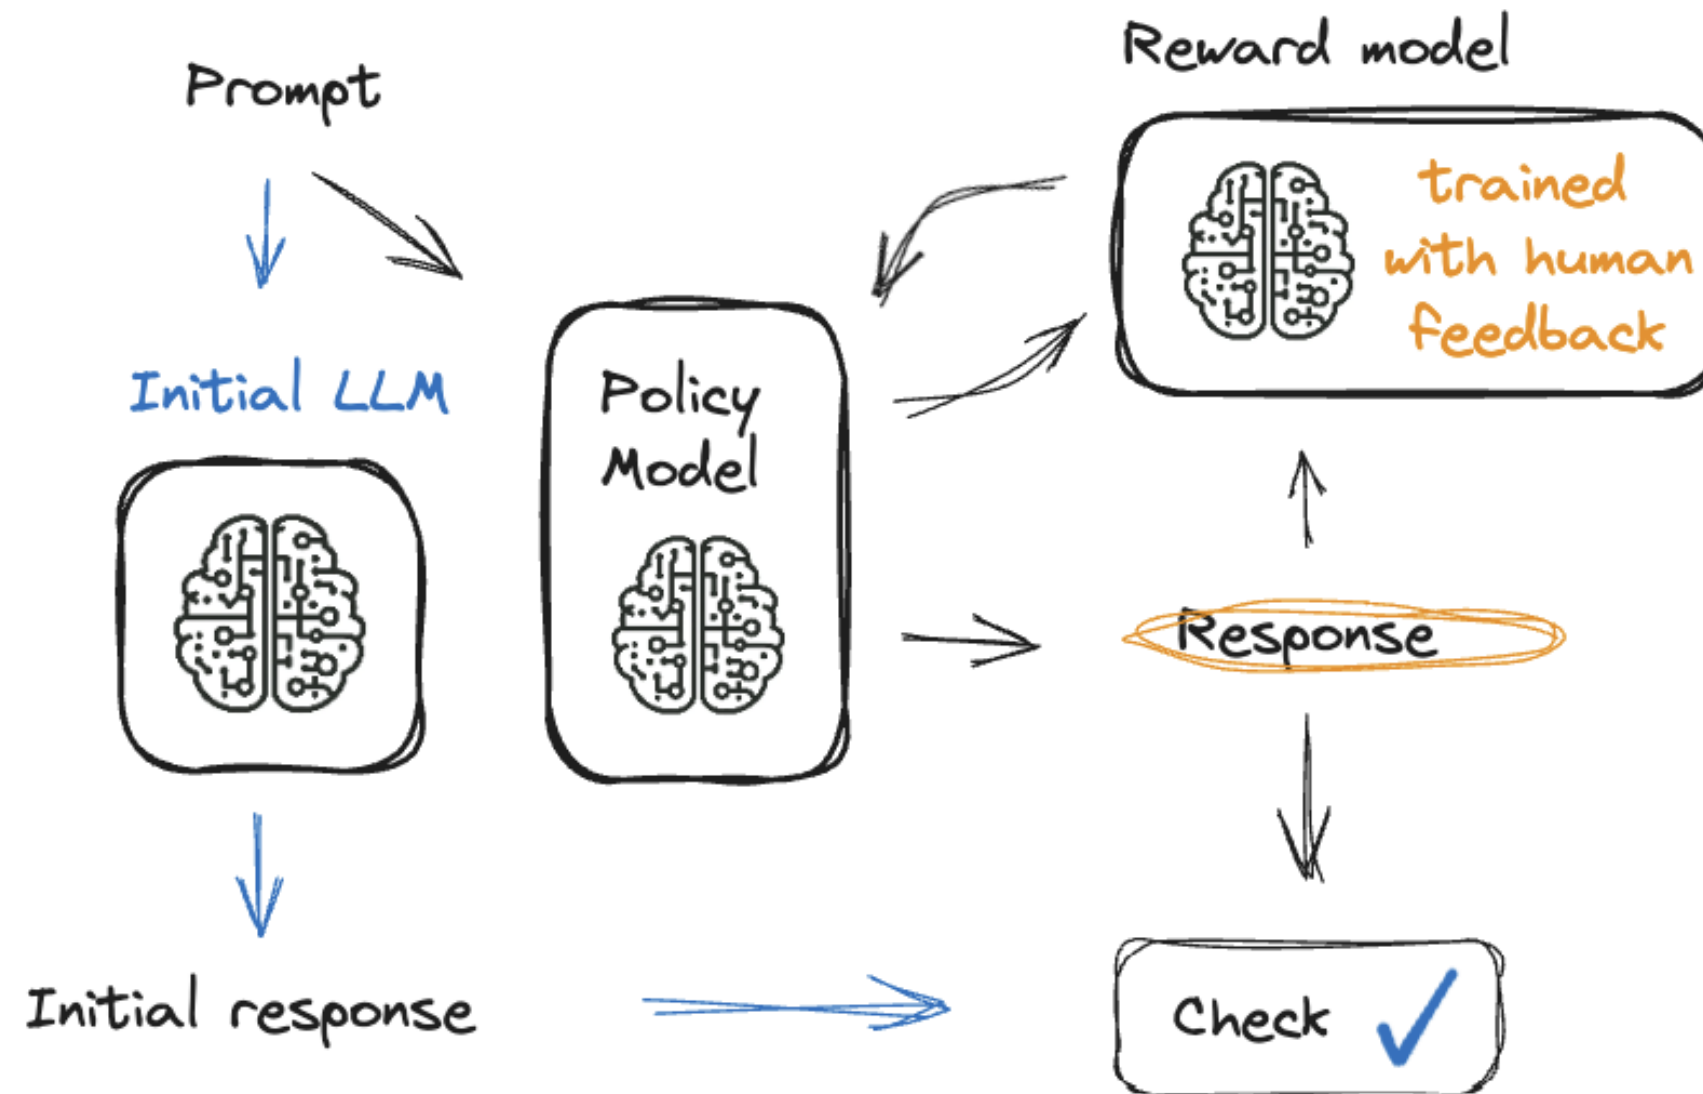
# Human in the loop systems



# Active learning in RLHF



# Active learning in RLHF




# Active learning

data in




# Active learning






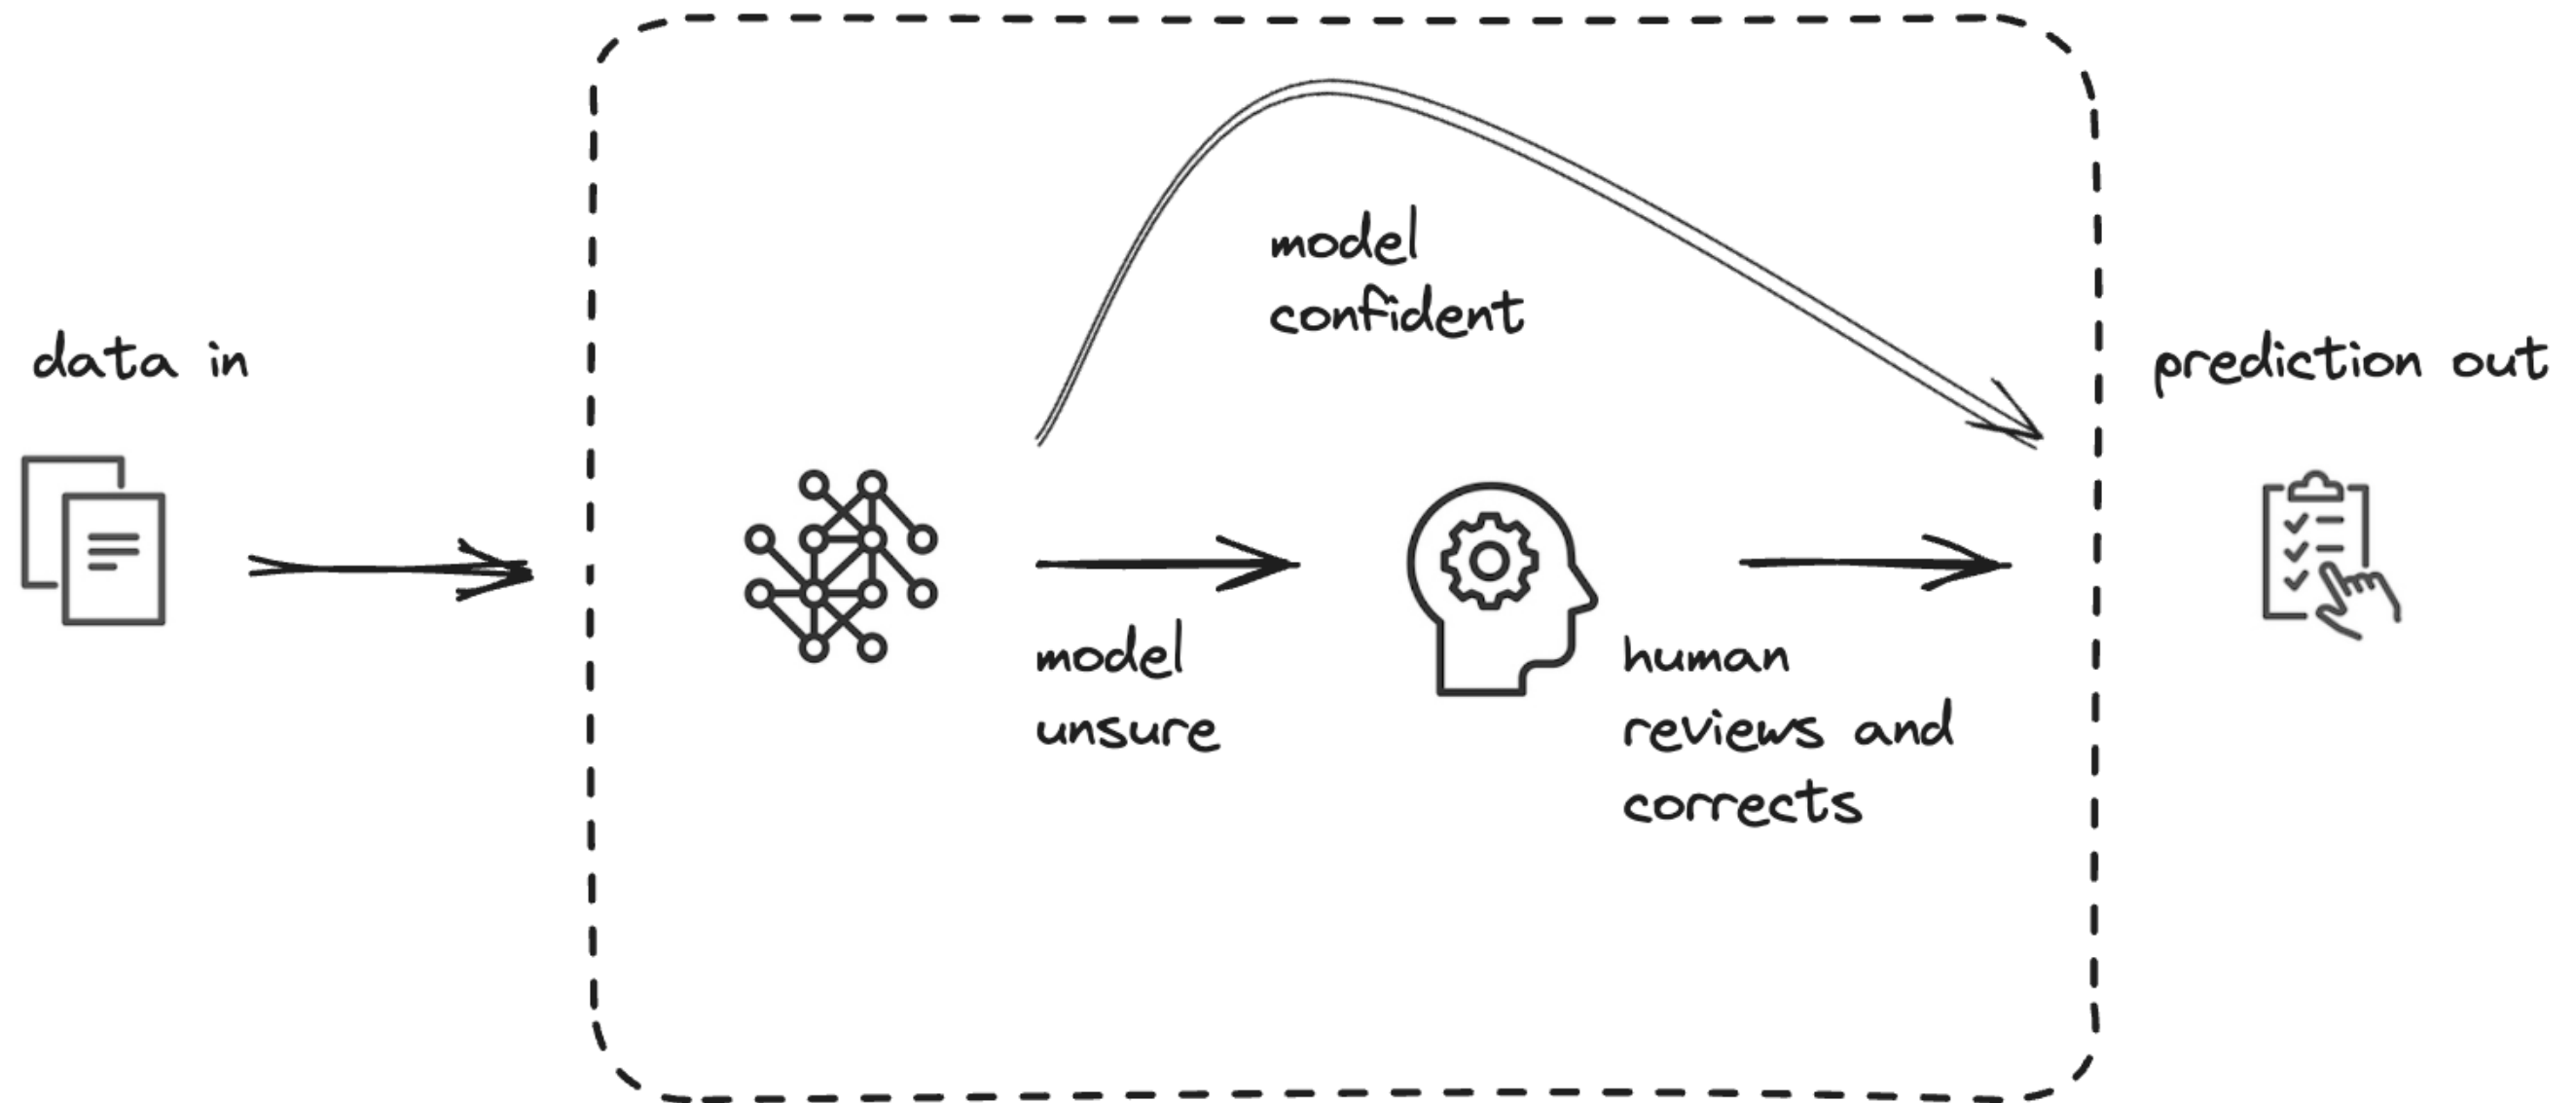
# Active learning



# Active learning



# Active learning



# Active learning pipeline with low confidence

```
from modAL.models import ActiveLearner
# Initialize learner
learner = ActiveLearner(
    estimator=LogisticRegression(),
    query_strategy=uncertainty_sampling,
    X_training=X_labeled, y_training=y_labeled
)
```

- **Uncertainty sampling:** points selected where confidence is lowest

# Active learning pipeline with low confidence

```
# Active learning loop
for _ in range(10):
    learner.teach(X_labeled, y_labeled)
    query_idx, _ = learner.query(X_unlabeled)
    X_labeled = np.vstack((X_labeled, X_unlabeled[query_idx]))
    y_labeled = np.append(y_labeled, y[query_idx])
    X_unlabeled = np.delete(X_unlabeled, query_idx, axis=0)
```

# Let's practice!

REINFORCEMENT LEARNING FROM HUMAN FEEDBACK (RLHF)